

The Shadow Files

Volume I
Foundations

YEHOR SELIN

Contents

Before You Begin	3
Who This Book Is For	5
How The Shadow Files Works	7
Prologue	10
What OSINT Really Is	12
The Rules	17
Protect Yourself First	31
Search Like a Pro	44
Username	57
Email	73
Phone Numbers	89
Your First Real Investigation	103
Epilogue	121
Appendix	123

Before You Begin

This book is an educational introduction to open-source intelligence — OSINT — written for beginners and for anyone who wants to understand how information about people and organisations can be found, verified, and protected using only what is already public.

Everything in these pages is taught inside a strict ethical and legal frame: open sources only, nothing that requires unauthorised access, and full respect for privacy and for the law. Laws differ from country to country — and more than one country's laws can apply to a single search. Know the rules where you are, and when you are not sure, get permission first.

Everything here is invented

Every name, username, email address, phone number, photograph, and case study in this book is entirely fictional. Each one was created to teach a method — never to describe a real person. Any resemblance to a real individual, account, or event is pure coincidence.

Why the examples are shown in full

So the lessons feel real, the examples use complete, ordinary-looking data instead of blanked-out symbols. Because every detail is invented, nothing here can touch a real person —

and you still see exactly how a real trail looks and behaves.

The rule that outranks every technique

Understanding how people become exposed exists here for one reason: to investigate responsibly, and to protect yourself and others better. Chapter 2 makes the ethical and legal boundaries concrete; for now, hold one line — use what you learn lawfully, ethically, and with care.

If you came here to protect yourself

Some readers open this book not to investigate anyone, but because they are worried about being found — by a stranger, an ex, a harasser. You are welcome here, and this book is for you too. You do not have to read eight chapters first: skip to Appendix D, "Take it back" for a plain, prioritised checklist of how to shrink your own footprint, then come back to the start whenever you are ready. Understanding how exposure works is exactly what lets you shut it down.

A professional is not the person who can find the most. It is the person who knows where to stop.

Who This Book Is For

This book was written for people starting from zero.

You do not need to be a hacker. You do not need to write code. You do not need special software or a powerful computer. If you can use a search bar, you already have everything you need to begin.

This is for you if you are

- Completely new and simply curious how investigators find things out.
- Someone who wants to protect yourself — to see what a stranger could learn about you, and shut those doors.
- A future analyst or investigator who wants a real, professional method from the very first page.
- A working professional — in HR, security, journalism, or law — who needs a structured, defensible way to check people and claims.

What you will be able to do by the end

Take one small clue — a username, an email, a phone number — and, step by step, responsibly build a clear and honest picture around it. Not by guessing. By method.

That is the whole promise of this first volume: the foundations
that everything else is built on.

How The Shadow Files Works

The series, this volume, and how each chapter is built.

The Shadow Files is a series of four volumes. Each one takes you one clear level higher, and each one assumes only what the last one taught.

- **Volume 1 — Foundations.** The mindset, the rules, and your first real skills. (*This book.*)
- **Volume 2 — Cross-Platform Attribution.** Following a person across social networks, messengers, photos, and leaks.
- **Volume 3 — Hard Targets.** Deleted accounts, geolocation, infrastructure, and the unusual tricks that surprise people.
- **Volume 4 — The Professional.** Corporate work, reports, counter-investigation, and turning skill into a craft.

A note on scope, so you know what you are holding: this volume is about finding and verifying people from the traces they leave. Hunting *infrastructure* — domains, servers, subdomains, exposed files, the recon a penetration tester lives on — is real OSINT too, but it is a different toolkit, and it waits for Volume 3. If you came for asset recon, you are early, not in the wrong place.

Inside this volume

1. What OSINT Really Is 2. The Rules: Law, Ethics, and the Lines You Never Cross 3. Protect Yourself First: OPSEC and the Sock Puppet 4. Search Like a Pro: Operators and Dorks 5. Usernames: The First Thread 6. Email: The Master Key 7. Phone Numbers: The Personal Anchor 8. Your First Real Investigation

How every chapter is built

The same simple rhythm, every time, so you always know where you are:

- The idea — in plain words, no jargon.
- The method — how a professional actually thinks and works, step by step.
- A worked case — a short, invented investigation you can follow move by move.
- An exercise — something safe and legal you do yourself.
- A recap — the three things to remember.

One promise about tools

This series teaches **method first, tools second**. Websites and apps come and go — some named here may be gone by the time you read this, and new ones will have taken their place. That is fine. The *thinking* never changes. Learn the method, and any tool becomes just a faster way to do something you already

understand.

Prologue

It always starts with one small thing.

A username someone chose years ago and forgot. A photo posted in a good mood. A phone number typed into the wrong box. A single line in an old profile.

None of it looks like much. That is exactly why it is so powerful.

People imagine that finding out about someone takes secret tools, or breaking in, or some dark corner of the internet. It doesn't. Almost everything that matters is sitting in the open, in plain sight, left there by ordinary people living ordinary lives. The work is not about breaking locks. It is about **noticing** — and knowing what to do with what you notice.

That is the craft this series will teach you. Not tricks. A way of seeing.

You will learn to take one small, forgotten thing and follow it — carefully, honestly — until it quietly tells you a story. You will also learn the other half, the half most people skip: how much *you* are giving away, and how to take it back.

I have spent years doing this work — in law enforcement, in private investigations, and teaching it to professionals around the world. And the most important thing I can tell you before we start is this: the power to see clearly comes with the duty to

use it well. We will talk about that duty often, because it is what separates a professional from everyone else.

Everyone leaves a shadow. By the end of this book, you will know how to read one — and, just as importantly, when to look away.

Every shadow has a name. Let's learn how to find it.

CHAPTER ONE

What OSINT Really Is

— *and how an investigator thinks*

Let's start with the honest version, no jargon.

OSINT means finding out about a person or a thing using only what's already out in the open. No hacking. No breaking in. No logging into anyone's account. Everything you'll ever do in this book uses doors that are *already unlocked* — public profiles, public posts, public records, public photos. You're not a burglar. You're someone who reads what everybody left lying on the table.

That's the first thing to unlearn: you don't need to be a "hacker." You need two things — the willingness to look, and a way to think. The looking is easy. This whole book is really about the thinking.

The one idea everything is built on

Here it is, the entire craft in one line:

Everyone leaves traces. You start from one trace, follow it to the next, and keep going until the pieces point to a real person.

We give those three moves names, and you'll hear them for the rest of the series:

- **1. Exposure** — the trace a person left behind (a username, a photo, a phone number, an old post).
- **2. Pivot** — using that trace to find the *next* one. One clue hands you another.
- **3. Attribution** — when enough clues line up and point at a single, real identity.

Exposure → pivot → attribution. That's it. Every case in every book is just this loop, run again and again.

Picture a loose thread on a sweater. You don't attack the whole sweater. You find *one* thread and pull — gently — and it leads to the next stitch, and the next. OSINT is pulling threads. The skill is knowing *which* thread, and *how hard* to pull.

How an investigator actually thinks

Four habits separate a professional from someone just typing a name into a search bar.

- **Start from one point.** Not "who is this person" — that's too big to answer. Instead: "I have this one username. What can it give me?" Small questions have answers.
- **Every answer is a new question.** Found an email? Good — now, *where else does that email live?* You are never finished; you just pivot to the next question.

- **Similar is not the same.** Two accounts using the nickname "shadowfox" are not automatically the same human. This is the mistake that ruins beginners. You don't *conclude* — you *confirm*, with a second and a third independent clue.
- **Watch your own wishes.** You will *want* the clue to point at your target — and that wanting is the danger. The moment you feel sure, deliberately hunt for the fact that would prove you *wrong*. That single habit will save you more grief than any tool.
- **It's structured curiosity.** Not magic, not luck. A calm, repeatable process anyone can learn.

A tiny example, so you can see the loop

Everything below is invented — a made-up person, to show the thinking, not the tools.

You're handed one thing: the username `j.karas_88`.

- **Step 1 — Exposure.** You check where that exact username exists across platforms. It turns up on a gaming forum and an old photo-sharing site.
- **Step 2 — Pivot.** The forum profile mentions a city, Riverton, and links a second nickname: `jaykaras`. A new thread to pull.
- **Step 3 — Pivot again.** That second nickname appears on a public profile that shows a first name, Jason, and a contact address: `jaykaras88@mailbox.com`.

- **Step 4 — Confirm, don't assume.** Same city on two sites. The same profile photo on both. The same odd way of writing ("cheers all!"). Three independent signals now point the same way.
- **Step 5 — Attribution.** You're no longer guessing. One username has become a named person — because each step *earned* the next.

Notice what did the work. No clever tool. The pivots did. Tools only speed up steps you already understand — which is exactly why this series teaches the method first.

The first time I did this for real, I started with a single username and a bad feeling. I did not expect much. An hour later I was looking at the whole shape of a person — where they worked, who they knew, the city they woke up in — built entirely from scraps they had left in the open, one small step at a time. Nobody handed me any of it; I had simply asked the right small questions in the right order. I have never forgotten that feeling. Not power, exactly — more like the quiet click of a lock you did not force, only understood.

The one rule that never bends

Everything here stays on the public side of a single line: open sources only, nothing you'd need permission for. The next chapter draws that line in full — where it runs, why it holds, and why staying on the right side of it makes you more effective, not

less.

Your first exercise

Investigate the easiest target you'll ever have — yourself.

- Write down the oldest username you ever used (that gaming or forum handle from years ago).
- Search it, exactly, across a few platforms.
- See what still exists, what it links to, and — honestly — how far a stranger could pivot from it.

You'll learn two things at once: how the loop *feels*, and how exposed you already are.

Recap

- OSINT is building a picture from **only public information** — no hacking.
- The whole craft is one loop: exposure → pivot → attribution.
- Method beats tools. Start small, pivot, and *confirm* before you conclude.

Next: the rules — where the line is, and why staying on the right side of it makes you more effective, not less.

CHAPTER TWO

The Rules

Law, ethics, and the lines you never cross

Start with the most reassuring thing I can tell you. Almost everything you will learn in this book is completely legal, completely open, and something you can do with a clear conscience. OSINT is built on looking at what people, companies, and governments have already made public. You are not breaking in anywhere. You are reading what is already on the shelf.

So this chapter is not here to scare you. It is here to hand you a compass. Once you understand a few simple ideas, you will *feel* where the line is, and you will not have to look it up every time. Here is the surprising part: the people who understand the rules best are also the boldest investigators. They move fast because they never worry that they crossed a line. Knowing where the wall is lets you run right up to it — so let me show you where it stands.

The bright line — looking versus touching

There is one line that matters more than all the others. Picture it as a bright, glowing line painted on the floor. Everything you do in OSINT falls on the safe side of it or the wrong side of it, and the line itself is simple.

Looking at what is already public is fine. Getting into what is locked is not.

That is the whole thing. In plainer words:

- *On the safe side:* reading a public profile, a company registration, a news article, a photo someone posted, a phone number listed in an open directory. All of this is out in the open. Anyone can see it. You are a member of the public looking at public things.
- *On the wrong side:* logging into someone else's account, guessing or using someone's password, sneaking past a login screen, getting into anything that was meant to be private and protected.

Think of a house on a street. You are allowed to walk down the street and look at the house. You can see the colour, the garden, the car in the driveway, the name on the mailbox. That is all public. You are *not* allowed to open the front door and walk in — even if you happen to know a key is under the mat. The moment you touch that door, you have crossed the line, no matter how easy it was.

In the computer world, the "front door" is anything with a password, a login, or a lock. If a step requires you to *sign in as someone else* or *get around a security control*, stop. That is the signal. That is the bright line, and you have reached it.

There are even laws with names attached to this idea. In the United States there is a law often called the *CFAA* — the Computer Fraud and Abuse Act. You do not need to memorise it. You just need its one-sentence heart: using a password you found to log into an account that is not yours is "unauthorised access," and it is a crime. Reading that an account *exists* is fine. Opening it is not. Same house, same door.

Public does not always mean fair game

Here is where beginners get tripped up.

"If it is public, I can do anything with it." That feels true. It is not. Something can be *out in the open* and still deserve *respect*: being allowed to *see* a thing is not the same as being allowed to *collect it, store it forever, and use it for any purpose you like*.

Public gives you permission to look. It does not hand you permission to do anything you want.

An example makes this obvious. Imagine someone posts a photo in a small support group for people recovering from an illness. Technically you can find it. But that person shared it in a small, gentle corner of the internet, expecting a small, gentle audience. If you rip it out of that corner and paste it into a report about them, you have done something that is *legal to see* but *wrong to use that way*. You broke the promise the setting made to them.

Professionals have a rule for this, and it is beautiful in its simplicity: respect the original context. Ask yourself, "Why did this person share this, and where?" The words may be public either way. The *spirit* is not.

There is also a whole family of privacy laws — in Europe the big one is called *GDPR* (it stands for General Data Protection Regulation, but the name does not matter). Again, one plain idea: even public information about a real, living person has to be handled for a good reason, kept only as long as you need it, and not hoarded "just in case." You do not collect a mountain of details about a person because you *can*. You collect what you actually need for the actual question you are answering. That is called *data minimisation*, and it is just good manners with a fancy name — take only what is on your list.

Consent, privacy, and the human on the other end

Behind every email address, every username, every blurry profile photo, there is a person. A real one. With a family, a bad day sometimes, feelings, a life. Keep that person in the room with you the entire time you work.

That single habit will keep you out of trouble more than any law will — because most wrong turns in OSINT do not come from not knowing the rules. They come from forgetting that the "target" is a human being.

So a few gentle standards:

- Have a *reason*. Before you look someone up, be able to say why. "I am curious" and "I want to build a file on my ex" are not reasons — they are red flags. A good reason is tied to a real, honest task.
- Do not *pretend to be someone else* to get information. Lying your way into data — telling a company you are the person's employer, for example — is called *pretexting*, and it is off-limits. Ask straight or do not ask.
- Do not *contact and trick* the person you are investigating. No fake friend requests to get past their privacy settings. No sweet-talking them into revealing things.

None of this makes you weak as an investigator. It makes you *trusted*. And in this field, being trusted is the whole career.

When you do this for work — extra rules apply

Everything so far applies to everyone. But the moment you run OSINT *for an organisation* — vetting a job candidate, checking an employee, screening a partner — a second layer of rules switches on, and breaking them can be its own kind of trouble even when every source was public. You do not need to be a lawyer. You do need to know the tripwires exist and loop in the people who handle them.

- **Hiring has its own laws.** In the United States, using background information to make a hiring decision can pull you under the *FCRA* (Fair Credit Reporting Act) — with

rules about permission and about telling the person before you act on what you found. And anti-discrimination law (the *EEOC* rules) means that if your search reveals someone's religion, health, pregnancy, age, or similar, and a decision then goes against them, the company can be accused of discriminating. The safe habit: record only job-relevant, verifiable facts — not everything you happened to see.

- Investigating your own people needs a documented reason. Regulated workplaces cannot just look up an employee "because." There has to be a legitimate, written purpose, and often HR or legal has to be in the loop *first*. Privacy laws like GDPR treat employee data with extra care. When in doubt, get authorisation before you start, not after.
- Faces and biometrics are special. Some places have specific laws about collecting or matching someone's face or other biometric data without consent. Treat face-search as a high-caution tool (we return to this later), and never run it on people at work without a legal check.

None of this is here to frighten you off. It is here so that the day your hobby skill gets used on the job, you already pause and ask, "Is there an extra rule because this is for an employer?" The answer is often yes.

The special case: breach and leak data

You are going to hear a lot about *breaches* and *leaks*, so let me define them plainly. A breach is when a company gets hacked

and its users' information spills out — sometimes including passwords. A leak is that spilled information floating around afterward. There are huge collections of this stuff on the internet.

This makes people nervous, and it should be handled with care. But there *is* a legal, useful, professional way to work with it — and there is a way that turns you into a criminal. The difference is, once again, one clean line.

You may check whether something was exposed. You may never use a stolen password to get in.

- **Checking exposure is fine.** There are respected services that let you type in an email address and learn, "Yes, this email showed up in the such-and-such breach in 2019." That is a *fact about exposure* — like reading that a store had a burglary. Useful, honest, allowed. (The best-known service even refuses to hand you the passwords on purpose; it only tells you *that* you were exposed.)
- **Using the stolen goods is not.** Taking a leaked password and typing it into someone's email to see if it still works is walking through the front door. It is unauthorised access. It is a crime. It does not matter that the password was "just sitting there."

Professionals sometimes use a leaked password in one very specific, clever, *safe* way: as a *fingerprint*. If the exact same

unusual password shows up attached to two different accounts, that is a strong hint the two accounts belong to the same person. Notice what happened there — the password was used to *connect the dots*, never to *unlock a door*. That is the entire game. Use identifiers to understand and link. Never to enter.

And one more honesty note about leak data: it is often old, messy, or flat-out wrong. A password from 2016 tells you about the past, not today. Never treat this stuff as gospel. Check it against a second source before you believe it.

How to stay on the right side — a simple method

Run this routine before every single lookup. It takes thirty seconds and it will keep you clean for your whole career.

- **Write down your reason.** One honest sentence: "I am checking this business before signing a contract with them." If you cannot write the sentence, do not do the lookup.
- **Ask: am I looking, or am I touching?** If any step needs a password, a login, or getting past a lock — stop. Full stop. Find another, open path or do not go there.
- **Ask: what did this person expect?** Respect the original context. A public post to the world is fair. A private corner is not, even if you can reach it.
- **Take only what you need.** Do not vacuum up a person's whole life. Collect the few facts that answer your actual question. Skip the rest.

- Never pretend, never trick, never contact under a false face. No lying to get data. No fake accounts to befriend the subject.
- Keep honest notes as you go. Save the link, save the date, save what you saw — exactly as you saw it. (More on this in a moment.)
- Check anything important twice. One source is a rumour. Two independent sources start to be a fact.

Run those seven checks and you will almost never put a foot wrong.

Keeping your notes honest

The notes deserve a moment of their own, because this is where good investigators separate themselves from sloppy ones — and it is also part of staying legal.

When you find something, write down *where* you found it, *when* you found it, and *what it actually said* — not what you wish it said, not your guess about what it means. Take a screenshot. Copy the exact web address. Note the date. Keep the fact and your *opinion about the fact* in two separate boxes.

Why does this matter so much? Two reasons.

First, honesty. If you blur "here is what I found" together with "here is what I think it means," you will fool yourself, and then you will fool whoever reads your work. Keep them apart. "The

record says the company was registered three months ago" is a fact. "This is probably a fake company" is your opinion. Different boxes.

Second, protection. If your work is ever questioned — and one day it might be — your honest, dated, sourced notes are your shield. They show you looked at public things, in a public way, for a real reason. Good notes are the difference between "I found this out properly" and "trust me." Always be able to prove the first one.

A short story: where a beginner slips and a pro stops

Here is a real-shaped situation. Everything and everyone in it is invented, but the choices are real. Follow the thinking.

Imagine you have been asked to check out a small vendor before your company signs a deal with them. The contact person is named Daniela Voss, and her email is daniela.voss@brightfield-supplies.com. Here is how it unfolds.

- **Step 1 — The open look.** You search the company name and find their website, their registration date, and a couple of news mentions. All public, all fair. You note the source and date for each. Good — you are firmly on the safe side.
- **Step 2 — The exposure check.** You run Daniela's email through a reputable breach-checking service. It tells you her email appeared in a leak two years ago. That is a *fact about*

exposure. You write it down: "email exposed in [leak name], 2024." Still safe. You only learned *that* she was exposed, not anything secret.

- **Step 3 — The temptation.** The same leak is floating around online, and it includes passwords. A beginner thinks: "I could just check if her old password still logs into her email. It would tell me so much." **This is the moment.** This is the front door with the key under the mat.

- **Step 4 — Where the pro stops.** A professional does not even reach for the door handle. Logging in with that password would be unauthorised access — a crime — and it would poison the entire job. Even if you found something, it would be useless and illegal. So you stop. You note the exposure as a *fact* and move on. That is it. That is the whole difference between the beginner and the pro: the pro feels the line and steps back without drama.

- **Step 5 — The clean finish.** You keep working the *open* trails instead. You check the company's registration against the official public register. You look for public reviews and any public court records. You corroborate anything worrying with a second source. You write a calm, honest report: facts in one column, your assessment in another, every claim dated and sourced.

You see how that went? At no point did staying legal make you *less* effective. It made your work *stronger*, because everything in your report can stand up to daylight.

The mindset: a pro knows where to stop

One idea ties this whole chapter together.

A beginner asks, "Can I get this?" A professional asks, "Should I, and how far do I go?"

Knowing where to stop is not a limit on your skill. It *is* the skill. Anyone can be reckless; it takes a professional to move powerfully through public information and then, at exactly the right moment, step back from the locked door. That restraint is what makes people trust you with serious work, and it is what lets you sleep well.

I still remember the first time a trail led somewhere I had no business going. The question I had been asked was small and fair. But the path kept opening — a home, a family, a daily routine — and I could feel how easy the next step would be. I closed the file. Not because I could not take that step, but because taking it would have turned an honest inquiry into surveillance of a life. Walking away felt like losing. It was the most professional thing I did all year.

You now have the compass. The rest of this book is about learning to walk. Let us begin.

Exercise

Here is a safe, legal exercise to make the bright line feel real.

- Pick a *large public company* you already know — a well-known brand, not a private person. Companies are public entities, so this keeps you far from any privacy concern.
- Spend fifteen minutes gathering only *public* facts: when it was founded, where it is headquartered, its official website, one or two recent news stories. For each item, write down the exact web address and the date you looked.
- Now write two short lists side by side. On the left: "Facts I found." On the right: "My opinion about what they mean." Keep them strictly separate. Feel how different those two columns are.
- Finally, write one honest sentence at the top of the page: your *reason* for looking. Practise the habit of never starting without it.

You just ran the whole method — reason first, public only, honest sourced notes, facts kept apart from opinion — on a target where you cannot possibly do harm. That is exactly how the pros warm up.

Recap

- The bright line is looking versus touching. Reading what is public is fine; logging into anything locked, or using a stolen password, is off-limits — always.

- Public does not mean anything-goes. Respect why and where something was shared, take only what you need, and remember there is a real person on the other end.
- A pro knows where to stop. Keep honest, dated, sourced notes, check important things twice, and treat stopping at the locked door as a skill — because it is the one that makes you trusted.

CHAPTER THREE

Protect Yourself First

OPSEC and the sock puppet

A small confession to start. The very first time I looked someone up online, I did it wrong.

I used my own account. My real name. My real photo. I opened the person's profile, poked around, and felt clever. Then, a few days later, that person showed up in my "People You May Know" list. The platform had quietly noticed I kept viewing them. It suggested us to each other. My cover — such as it was — was gone before I even knew I had a cover to lose.

Nobody got hurt that time. But it taught me the most important lesson in this whole book, and I want to give it to you before anything else.

Before you look at anyone, protect yourself.

Investigation is the second step. Safety is the first.

This chapter is about OPSEC. That word comes from the military. It means *operational security* — a fancy way of saying "don't leave a trail that points back to you." For us, it means one simple thing: you can look at the open internet without the open internet looking back at you.

Every idea here stays plain, and the setup is one you can build on a normal laptop this afternoon. No spy gadgets. No money required beyond a coffee's worth.

Why you never investigate from your own accounts

The core idea. Say it out loud if it helps.

Your real accounts are wired to your real life. The moment you use them to look, you risk showing your face to the person you're looking at.

Modern platforms — Facebook, Instagram, LinkedIn, TikTok, all of them — are watching *your* behaviour as much as everyone else's. When you view a profile again and again, the platform notices. It may suggest you two to each other. On LinkedIn, the person can often see *exactly* who viewed them. One careless click and your name lands in their notifications.

There's a second danger, quieter but just as real. Your own account is full of *you*. Your friends. Your photos. Your workplace. Your habits. If the person you're researching is even a little curious, or a little dangerous, your account is a map straight to your front door. In OSINT, the person you study can sometimes study you back. Never hand them the map.

So we build a wall between "you the person" and "you the investigator." That wall is called a sock puppet.

What a sock puppet is (and what it is not)

A sock puppet is a clean, research-only account. Think of it as a work glove. You put it on to handle things you don't want to touch with your bare hand. It has no link to your real name, your real email, your real photo, or your real phone. It exists for one job: looking, quietly.

The name comes from the old trick of putting a sock on your hand and pretending it's a different character. Same idea. It's a separate face you wear for research.

Beginners get into trouble on what a sock puppet is **not**, so be clear on it.

- It is **not** for tricking people into talking to you. We don't message the subject, friend them, or join their private groups uninvited — that crosses from *looking* into *interacting*, the line Chapter 2 drew.
- It is **not** a fake identity for scamming, catfishing, or harassment. Ever. If that's the goal, close this book.
- It is **not** a magic cloak. It hides your name. It does not, by itself, hide your location or your computer. We'll fix those separately.

The best sock puppets follow what missing-persons investigators call the **no-touch rule**: observe only. No likes. No comments. No friend requests. No password-reset probes. You are a fly on a public wall, nothing more.

A sock puppet is a research-only account for watching the open web — never for contacting, tricking, or engaging anyone.

Building your safe setup, step by step

Let's build the whole thing, in order. Do it once, keep it, and you'll have a clean research kit forever.

- **Step 1 — Make a separate browser profile.** You don't need a second computer. Most browsers (Firefox, Chrome, Brave) let you create a fresh profile — a totally separate browser space with its own history, cookies, and logins. Create one and name it something dull like "Research." Never open your personal accounts inside it, and never open your research accounts outside it. This one habit does most of the heavy lifting.
- **Step 2 — Harden that browser a little.** In the research profile, turn off telemetry (the browser phoning home about what you do). Block third-party cookies. Install a good ad-and-tracker blocker like **uBlock Origin**. These are free, take five minutes, and stop most of the invisible tracking that follows people around the web. (Tool names change over the years — the *idea* is "block trackers, kill cross-site cookies.") Pick whatever reputable tool does that when you read this.)
- **Step 3 — Understand what a browser leaks.** Even with a blocker on, your browser quietly shows a lot: your operating system, your screen size, your fonts, your graphics chip.

Together these make a **fingerprint** — a pattern almost as unique as a real fingerprint. A VPN does *not* hide this (more on VPNs in a moment). You don't need to become an expert today. Just know it exists, and keep all your research in that one browser profile so your "investigator self" always looks the same. Consistency is safer than chaos here.

- **Step 4 — Create a brand-new email.** Sign up for a fresh mailbox (Gmail, Proton, Outlook — any reputable one). Use it for *nothing* but research. This email must never have touched your real name, your real recovery phone, or your real accounts. It is the root of your sock puppet. Treat it like the key to the whole thing, because it is.
- **Step 5 — Design the persona on paper first.** Before you make any social account, decide who this puppet *is*. A plausible name. An age. A city. A boring job. A couple of harmless interests. Write it down and keep it consistent. A tool that generates random-but-realistic identities can seed this for you. The goal is *forgettable and ordinary* — someone nobody looks at twice.
- **Step 6 — Handle phone verification honestly.** Some platforms demand a phone number to sign up. Do not use your personal number — that instantly links the puppet to you. If you'll do this regularly, a cheap dedicated research SIM (a burner number kept only for this) is the clean answer. Free internet-phone numbers often get rejected, so a real SIM saves headaches.

- **Step 7 — Pick a profile picture carefully.** Never use a stock photo or a celebrity — a quick reverse-image search would burn you in seconds. Many people use an AI-generated face, but crop it and strip the hidden data (metadata) from the file first, because that data can quietly reveal where and how the image was made. Honestly, for pure passive watching, no photo at all is often fine.
- **Step 8 — Age the account before you use it.** A brand-new account is a red flag to platforms and people alike. Log in over a few days. Follow a few neutral pages — a sports team, a news site, a cooking group. Let it look lived-in. Then, and only then, go near your actual research. New accounts get locked fast; a warmed-up one survives.
- **Step 9 — Never mix things.** This is the golden thread through all of it. One email per puppet. One phone per puppet. Never reuse a photo, a password, or even a favourite turn of phrase across puppets or with your real self. Reuse is how investigators get *linked* — to each other and back to a real name. Keep everything in separate lanes.

That's the whole kit: separate browser profile, blocked trackers, fresh email, a thought-out persona, a clean phone if needed, a careful photo, and a warmed-up account. On a normal laptop. Today.

What a VPN does — and the honest debate about it

You'll hear the word VPN everywhere in this world, so here it is, simply.

A VPN (virtual private network) hides your real internet address — your IP address, which is roughly the return address your connection stamps on everything you do. Without a VPN, the site you visit can see an address that points toward your home town and your internet provider. With a VPN, the site sees the VPN's address instead. It's like mailing a letter through a forwarding office so nobody reads your home street.

A few plain rules:

- **Use a paid, reputable, no-logs VPN.** "No-logs" means they don't keep records of what you did. Free VPNs make their money by selling your data — exactly the thing you're trying to protect. Skip them.
- **A VPN hides your address, not your fingerprint.** It is one layer, not the whole wall. This is why we bothered with the browser profile.

Now the honest part. Experienced investigators *disagree* about VPNs, and I want you to hear both sides rather than a slogan.

One camp says: always use a VPN. A clean address alone isn't enough against fingerprinting and location clues, so keep it on for everything.

The other camp says: turn the VPN off for the one moment you create an account. Why? Because some platforms recognise VPN addresses and treat them as suspicious *at signup* — they throw up puzzles, demand a phone number, or ban the new account outright.

Both camps are right, in context. Here's how I reconcile it in plain terms:

Use your VPN for everyday looking and collecting. But for the delicate moment of creating a sock puppet on a picky platform, a clean, ordinary-looking home connection often works better than a VPN address the platform distrusts.

Don't memorise a rule. Understand the tradeoff: a VPN protects you, but a *flagged* VPN address can get your new account killed on sight. Judge each moment.

Saving evidence honestly

Finding something is only half the job. If you can't *prove* what you found, and prove you didn't tamper with it, it may as well not exist. This matters double if your work could ever end up in a report, a dispute, or a court.

Here's the beginner-friendly version of doing it right:

- Capture the whole page, not a tiny crop. Save the full page with the web address (URL) visible and the date on screen. A

lonely cropped screenshot with no context is easy to dismiss.

- **Write down when and where.** The exact date and time you saw it, the URL, and how you got there (what you searched, which platform). "I found it somewhere on Facebook" is worthless. "Viewed at this URL on this date at this time" is solid.
- **Save it, don't just look at it.** Content gets deleted. Save the page, and consider using a public web archive so there's an independent, dated copy you didn't control.
- **Keep the original untouched.** Don't crop, brighten, or "clean up" the saved evidence. If you need an edited version for a report, keep the untouched original too, and note that you edited a copy.
- **Save the misses as well as the hits.** If you searched and found *nothing*, record that too. It proves you looked, and it stops you from re-treading old ground.

As you grow, there are proper tools that capture pages automatically and stamp each one with a tamper-proof seal and a trusted timestamp — so nobody can accuse you of quietly changing a screenshot afterward. You don't need them on day one. You *do* need the habits above from your very first search.

Simple note-taking so you never lose the thread

An investigation is a chain of small discoveries, and it is astonishingly easy to lose your place. Notes are how you hold

the thread.

I learned this the hard way — by not doing it. I was once deep into a case with a conclusion I was sure of and no memory of how I reached it. I could not retrace a single step, so I had to throw the whole thing out and start over. A finding you cannot walk someone else back through is not a finding — it is a feeling. Since that day I have written down every click, in order, before I let myself believe anything.

Keep it dead simple. One note file per case. For every finding, jot four things: what you found, where (the URL or source), when you found it, and how (the search or tool). And keep one line sacred:

Separate what you saw from what you think it means. Facts in one column, your guesses in another. Never let them blur.

That last habit is the difference between a professional and someone who fools themselves. "His profile lists Kyiv" is a fact. "So he probably lives in Kyiv" is a guess. Write both, but never confuse them. Also record your dead ends and the reasons you dropped a lead — future-you will thank present-you.

A beginner slip, and the fix

Here is a short, made-up story that could happen to anyone.

Meet Daria Kovalenko, a beginner, curious and eager. A friend asks her to check whether a marketplace seller named Roman Petrenko is trustworthy — a small favour, or so it seems.

- **Step 1 — The slip.** Daria opens Instagram in her normal browser, logged in as her real self, and views Roman's profile. Twice. She also clicks a link he posted, which loads a page that quietly records visitors. She feels productive.
- **Step 2 — The exposure.** A day later, Roman's account shows up as a suggested friend for Daria — and hers for him. Because she was logged in as herself, her real name and photo were the "visitor" the platform tied to those profile views. Roman now knows a woman named Daria Kovalenko has been looking at him. Her real face. Her real friends list, one click away.
- **Step 3 — The fix, part one.** Daria stops. She creates a separate "Research" browser profile and never touches her personal accounts inside it. She installs a tracker blocker. From now on, her looking and her real life live in different rooms.
- **Step 4 — The fix, part two.** She makes a fresh email, invents a plain, boring persona named Olena Marchuk, and builds a lightweight sock-puppet account. She warms it up for a few days with neutral follows. She keeps her VPN on for browsing.

- **Step 5 — The fix, part three.** Now she views Roman's public profile through the puppet, in the research browser. She saves the full page with the URL and date, notes what she found versus what she suspects, and — crucially — never messages him. Observe only.

Same task. Same laptop. The difference is a wall between Daria-the-person and the eyes doing the looking. That wall is the whole chapter.

Exercise

Build your kit before you ever need it. This should take under an hour and costs nothing but a fresh email.

1. Create a separate browser profile just for research. Name it plainly.
2. In it, install a reputable tracker blocker and turn off telemetry and third-party cookies.
3. Create a brand-new email that has never touched your real identity.
4. On paper, design one ordinary persona: name, age, city, a boring job, two harmless interests. Keep it forgettable.
5. Visit a "what does my browser reveal" test page and simply *read* what your research browser leaks. Don't fix anything yet — just look, so the idea of a fingerprint feels real.
6. Open a plain note file and practise logging one harmless, public thing (say, a museum's opening hours) using the four fields: **what, where, when, how** — and add one line separating fact from guess.

You've now built the wall and the notebook. Everything else in this book stands on top of them.

Recap

- **Safety comes first, always.** Never investigate from your own accounts — build a wall between you-the-person and you-the-investigator using a separate browser profile and a research-only sock puppet.
- **A sock puppet is for watching, not touching.** Fresh email, plain persona, clean phone if needed, careful photo, warmed-up account — and never reuse anything across puppets or with your real self.
- **Work honestly and keep the thread.** Use a reputable VPN (but understand the signup debate), save full-page evidence with URL and date, keep originals untouched, and take simple notes that separate what you *saw* from what you *think*.

CHAPTER FOUR

Search Like a Pro

Operators and dorks

Most people search badly. That is not an insult. It is just the truth, and I include my younger self in it. We type a few words into a search box, glance at the first page, and give up if the answer is not staring back at us. The search engine did exactly what we asked. We simply asked a vague question, so we got a vague answer.

This chapter fixes that. By the end of it you will search on purpose. You will tell the engine precisely what you want and, just as importantly, what you do *not* want. You will learn a small set of commands — we call them **operators** — that turn a wide, messy search into a narrow, sharp one. And you will learn "dorking," a word that sounds mysterious and is not. It just means searching carefully.

Everything here uses public search engines the way they were built to be used. Just skill.

The one big idea: precision beats volume

Almost everyone makes the same mistake. They think more results are better — that a search returning two million pages is a success.

It is the opposite. Two million results means the engine did not understand you. It threw everything vaguely related into a pile and left you to dig. A great search often returns *five* results — and the answer is in the first one.

Your goal is not more results. Your goal is the right results.

It works like fishing. A beginner drags a huge net through the whole ocean and hauls up mostly seaweed. A professional knows the exact spot, the exact bait, the exact depth — and catches the one fish they came for. Operators are your bait and your depth gauge.

One key term first, because I will use it constantly. An operator is a special word or symbol you add to your search that gives the engine an extra instruction. Normal words say *what* to look for. Operators say *how* and *where* to look. That is the whole secret.

The operators, one at a time

Each one comes with a tiny example. Read them slowly and, if you can, type each into a search engine as you go. You learn this in your fingers, not just your eyes. Throughout, I use made-up names, handles, and websites so nothing here points at a real person.

Exact match — quotation marks

Put words inside "quotation marks" and the engine must find those exact words, in that exact order. Without quotes, the engine feels free to shuffle your words around, swap in synonyms, and drop words it thinks are unimportant.

- Marina Delcroix — the engine may show pages with "Marina," pages with "Delcroix," and pages that only loosely match.
- "Marina Delcroix" — only pages with that full name, spelled that way, side by side.

This is the single most useful habit in searching. Names, phrases, error messages, quotes from a speech — wrap them in quotes and watch the noise fall away.

Search inside one website — site:

The site: operator tells the engine to look *only* within one website. You type site: then the web address, no space.

- site:linkedin.com "Marina Delcroix" — find that name only on LinkedIn.
- site:reddit.com stolen bicycle Copenhagen — search Reddit alone.

This is powerful for two reasons. First, a website's own search box is often weak, while a search engine has already indexed that whole site far better. Second, you can aim at exactly the corner of the internet where your answer probably lives.

You can even point *site:* at a whole *type* of website by using the ending of the address:

- *site:gov.uk* — only British government pages.
- *site:.edu* — only educational institutions, mostly universities. (Note the leading dot: *site:.edu* matches the whole *.edu* family, not a single site.)

Find specific file types — *filetype:*

Not everything on the internet is a normal web page. Reports, spreadsheets, presentations, and forms are often documents. The *filetype:* operator finds them.

- *filetype:pdf "annual report" Meridian Logistics* — find PDF reports for a made-up company.
- *filetype:xlsx budget 2024 site:gov.uk* — find spreadsheets (*xlsx* is an Excel file) on UK government sites.
- *filetype:pptx "Meridian Logistics"* — find PowerPoint slide decks that name the company.

Documents are gold. People are careful about what they put on a web page but often forget what is buried in an old PDF or a spreadsheet that got posted years ago. A phone list, a staff directory, meeting minutes — these hide in documents.

Look in the title or the web address — *intitle:* and *inurl:*

Every web page has a *title* (the headline that shows in the browser tab) and an *address* (the URL, the *www...* line at the

top). Two operators let you search inside just those parts.

- `intitle:"staff directory" Meridian Logistics` — pages whose *title* contains "staff directory." A page bothering to title itself that is very likely the real thing, not a page that merely mentions it in passing.
- `inurl:careers site:meridianlogistics.example` — pages whose *address* contains the word "careers."

Why does this matter? Because a keyword in the title or the address is a strong signal that the page is *about* that thing, not just brushing past it. It is a way of finding pages that are on-target rather than tangentially related.

Remove what you do not want — the minus sign

Put a minus sign - directly in front of a word (no space) and the engine will throw out every page containing that word. This is how you clean up a messy search.

- `jaguar -car -animal` — you want the guitar amplifier or the sports team, not the car and not the big cat.
- `"Marina Delcroix" -obituary` — you are looking for the living person and keep hitting a funeral notice for someone with the same name. Remove it.
- `Meridian Logistics -jobs -careers` — you want news and records about the company, not endless job postings.

The minus sign is quiet but mighty. Most bad searches are not missing the right pages — they are drowning them under the wrong ones. Subtract the wrong ones.

Widen the net on purpose — OR

Sometimes you want *either* of two things. Type OR in capital letters between them. (Capitals matter; lowercase "or" is treated as a normal word.)

- "Marina Delcroix" OR "M. Delcroix" — catch both the full name and the shortened form.
- Meridian Logistics (invoice OR receipt OR statement) — the round brackets group the choices neatly.

Use OR when your target might be written more than one way — a nickname, an abbreviation, a maiden name, a spelling variant. You are telling the engine, "any of these will do."

Fill in the blank — the wildcard

The asterisk * is a wildcard. It means "something goes here, I do not know exactly what." It works best inside a quoted phrase.

- "Marina * Delcroix" — finds "Marina Anne Delcroix," "Marina J. Delcroix," and so on. The * stands in for the missing middle name.
- "chief * officer" site:meridianlogistics.example — finds "chief executive officer," "chief financial officer," "chief technology officer," all at once.

The wildcard is your friend when you know the shape of a phrase but not every word in it.

Narrow by time — date filters

Old results can bury fresh ones, and sometimes you specifically want *old* ones. Every major search engine has a date filter, usually tucked under a "Tools" or "Filters" button beneath the search box. Set it to the last year, the last month, or a custom range.

On some engines you can also filter by date right in the query with *before:* and *after:* using year-month-day:

- Meridian Logistics *after:2023-01-01 before:2023-12-31* — only results from the year 2023.

Dates matter more than beginners expect. A merger, a name change, a scandal — knowing *when* something happened lets you slice straight to the reporting from that window.

Use more than one search engine

One habit separates the pros from everyone else: they never trust a single search engine.

Google, Bing, and Yandex do not show the same results. Each one crawls the web differently, ranks pages differently, and — crucially — censors and prioritises differently based on where its company sits.

- Google is the broadest for most of the world and the one most people know.
- Bing (Microsoft's engine) often surfaces pages Google buries, and its image search behaves differently, which sometimes helps.
- Yandex is a Russian engine and is far stronger for content in Russian and for many things in Eastern Europe and Central Asia. It is also famously good at reverse image search — but that is a later chapter.

Different engines see different internets. If one shows you nothing, another may show you everything.

The good news: the core operators — quotes, site:, filetype:, the minus sign — work across all three, with only small differences. So the skill you just learned travels. When a search comes up empty on Google, do not conclude the answer is not out there. Run the same query on Bing and Yandex before you give up. I have cracked more than one case simply by moving the exact same search to a different engine.

So what is "Google dorking"?

You have probably heard the phrase "Google dorking" and maybe it sounded like hacking. It is not.

A "dork" is just a precise search that stacks several operators together. That is all. The word came from security researchers,

but the technique is nothing more than what you have already learned in this chapter, combined. When someone writes a clever query like `site:meridianlogistics.example filetype:pdf "employee handbook"`, that is "a dork." It is skilled searching, not sneaking.

Finding a document that a website put on the open internet is legal and normal — a search engine indexing a public file makes reading it no different from reading any other public page. The line Chapter 2 drew still holds: the operator is neutral, and it is your intent and actions that matter. If a search accidentally surfaces something clearly sensitive that was exposed by error, the right move is to note it, not exploit it, and where appropriate report it.

A dork is a precise question. Asking a precise question is never the crime. What you do with the answer can be.

A worked case: finding one person, the professional way

Watch the operators combine into something real. A fictional request: a colleague asks me to confirm the current employer and professional background of a person they only know as Marina Delcroix, who they believe works somewhere in the logistics industry in Europe. A name and a hunch — that is all we have. Watch how each operator narrows the fog.

- Step 1 — Lock the name. I never search a name without quotes. I start with "Marina Delcroix" in Google. A pile of loosely related results comes back, including, annoyingly, an obituary for a different Marina Delcroix and dozens of unrelated social posts. That pile tells me the name exists and is not vanishingly rare. Good enough to begin.

"Marina Delcroix"

- 1 Obituary: Marina Delcroix, 1948-2022
 - 2 Marina Delcroix (@delcroix_art) photos
 - 3 Delcroix family reunion – guestbook
 - 4 Marina Delcroix logistics – LinkedIn
- 3 of top 4 are noise; only #4 is on-topic

- Step 2 — Cut the noise. I remove the wrong person and the wrong context: "Marina Delcroix" -obituary -funeral. The death notice vanishes. The results are cleaner already. Precision is mostly subtraction.

- Step 3 — Go where professionals live. Logistics is a career, so I aim at the professional network: `site:linkedin.com` "Marina Delcroix" logistics. Now I get a small handful of profiles instead of the whole internet. The top profile looks perfect — a Marina Delcroix, and the word "logistics" is right there. It would be easy to stop and call it a match. But I read past the headline, and the friction shows: this Marina Delcroix does *supply-chain software sales*, based in Montreal, and "logistics" is just a keyword in her pitch. Wrong person. The engine matched my words, not my target. A first-page result that fits your query is not the same as a result that fits

your case.

- **Step 4 — Reframe, then confirm with a second source.** So I stop trusting the top hit and tighten the query with the one detail I actually have — Europe: site:linkedin.com "Marina Delcroix" logistics (Europe OR Rotterdam OR Antwerp).

This time a different profile rises: a Marina Delcroix in operations at a firm I will call Meridian Logistics, based in the right part of the world. That is a lead, not a fact — so I corroborate it before I believe it. I never trust one page, so I search the company's own site for her name:

site:meridianlogistics.example "Marina Delcroix". A team page and a press release both name her as a regional operations manager. Now an independent source agrees with the profile, and the Montreal false hit is behind me.

- **Step 5 — Reach for the documents.** Web pages get scrubbed; documents get forgotten. I run "Marina Delcroix" Meridian filetype:pdf. A conference programme from 2023 lists her as a speaker, with a short professional biography — exactly the background detail my colleague asked for. That biography also mentions a *previous* employer, which is a fresh thread I could pull if I needed to.

- **Step 6 — Check a second engine.** Before I write anything up, I run the two best queries again on Bing and Yandex. Bing surfaces a local news article Google had buried, which quotes her by name and title. Now three separate, public sources agree on employer and role. I stop here, because the question

is answered and everything traces back to a public page I can cite.

Notice what happened. I never guessed. Each step used a plain operator you now know, and each step either narrowed the search or added an independent confirmation. That is the whole craft in miniature.

And notice the deeper lesson: good searching is itself the first pivot. In OSINT, a "pivot" means using one fact you have found to reach the next fact. A precise search does not just answer your question — it hands you new threads (a previous employer, a co-author, an old address) that become the starting point of your next search. You are not just finding an answer. You are opening doors.

Exercise

Pick a made-up target you can safely play with, or use a public figure you have permission and reason to look up for practice — a well-known author, for instance. Then do the following, writing down the query and the number of results at each step so you can *see* precision at work:

1. Search their full name with no quotes. Note roughly how many results come back.
2. Search the same name inside quotes. Did the number drop? Did the top results get more relevant?
3. Add a site: operator to search only one platform where they are likely to appear.
4. Add a minus sign to remove one kind of

result that is getting in your way. 5. Run a `filetype:pdf` search for their name and see whether any documents surface that the normal web pages did not show. 6. Finally, take your single best query and run it, unchanged, on **Bing and Yandex** as well as Google. Write down one result that appeared on one engine but not the others.

If you finish and your last search returns fewer, sharper results than your first — congratulations. You have stopped dragging a net through the ocean.

Recap

- **Precision beats volume.** A search that returns five right pages beats one that returns two million wrong ones. Operators — quotes, `site:`, `filetype:`, `intitle:/inurl:`, the minus sign, **OR**, the wildcard `*`, and date filters — are how you get there.
- **Stack operators and switch engines.** A "dork" is nothing but several operators combined into one precise, legal question. And because Google, Bing, and Yandex see different internets, running the same query on all three is often the difference between empty-handed and answered.
- **Good searching is the first pivot.** Each sharp search not only answers today's question but hands you the next thread to pull — always confirmed against a second independent public source before you believe it.

CHAPTER FIVE

Username

The first thread

Here is the day a username unravelled a whole life.

I once had almost nothing to work with. One handle. Seven little letters that a person had typed into a signup box years ago and forgotten about. No real name. No photo I trusted. No location. Just that string.

Two hours later I had a full name, three email addresses, a home city, and a face. I did not hack anything. I did not break a single law. I just pulled on one thread, and the whole sweater came apart.

That thread was a username. And in this chapter I am going to teach you how to pull it.

What a username actually is

A username is the name a person picks for themselves online. It is also called a handle. It is the @something on Twitter or Instagram. It is the name in front of a Reddit account. It is the part of an email before the @ sign.

People choose these names once, when they are young or in a hurry, and then they keep them. Forever. That is the secret. A

username feels small and throwaway. It is not. It is one of the strongest starting clues in all of open-source intelligence.

Open-source intelligence, or OSINT, just means finding things out using information that is already public — only what people have left in the open, where anyone is allowed to look (the boundary Chapter 2 set). We collect it, we connect it, and we reason carefully about what it means.

A username is the perfect first thread for one plain reason.

People are lazy with names. They reuse the same handle everywhere.

Look at your own life. You probably have one or two nicknames you have typed into a hundred signup forms. Maybe it is a childhood nickname. Maybe it is your name plus a favourite number. Maybe it is a word you liked when you were fourteen. And you have used it on gaming sites, forums, shopping accounts, and social media, over and over, for years.

Everyone does this. It is human. Making up a fresh clever name every single time is tiring, and people are tired. So they reuse. And that reuse is a trail.

The core idea

The whole chapter in one sentence:

If a person uses the same rare username in many places, you can walk from one place to the next and slowly build a picture of who they are.

That is it. One handle becomes a map. Each stop on the map gives you something new: a photo, a real name, a linked email, another handle, a city. And each new thing lets you take another step.

But there is a rule I need you to burn into your mind before we go one inch further. It is the most important sentence in this book.

Matching names do NOT prove same person.

Two accounts with the same handle might be one human. Or they might be two strangers who happened to pick the same word. Or a name someone abandoned and a stranger later grabbed. The match is a lead, not a conclusion. A lead is a hint that points somewhere. A conclusion is something you have actually proven. Never confuse the two. This is the mistake that ruins beginners, and it is the single sentence to carry out of this chapter.

I learned it the embarrassing way. Early on, I was certain I had my man. Same rare handle, same city, the timeline fit. I wrote his name at the top of my notes and felt clever. I was wrong. It was a stranger who happened to share the handle — the trail had forked two profiles earlier, and I walked past it because I wanted

to be right. That miss taught me more than any clean hit ever did: the moment you start hoping a clue points at your target, you have stopped investigating and started guessing.

The confidence ladder

Since matching names do not prove identity, we need a way to talk about *how sure* we are. I use a simple ladder. Every clue you find sits on one of these rungs.

- **Weak.** A single shared handle. A common avatar. A generic bio. On its own, this proves almost nothing. Especially if the handle is short or ordinary, like john or mike88.
- **Moderate.** A rare, unusual handle reused, plus something else that fits. Same account age. Some overlapping friends. Now you are getting somewhere.
- **Strong.** A truly rare signal lines up. The exact same unusual profile photo file. The exact same odd phrase in two bios. A recovery hint that matches. Two of these together and you are almost there.
- **Confirmed.** The subject links the accounts themselves ("find me on X"). Or a hard identifier, like a recovery email fragment, points to one and only one person.

Your job as an investigator is to climb this ladder. You start at the bottom with a weak lead. You gather more clues. Each clue that fits pulls you up a rung. You do not name a real human until you are near the top.

The rarer the clue, the faster you climb. A common handle barely moves you. A one-of-a-kind handle plus a one-of-a-kind photo can carry you most of the way in two steps.

The method, step by step

This is how I work a username from scratch. Follow the order.

- **Get your seed.** The one handle you start with is your seed. Write it down exactly, letter for letter. Case usually does not matter, but spelling always does.
- **Check where it exists.** Take the seed and look for it across many platforms at once. This is called **enumeration** — a big word that just means "checking a list." You are asking: on which sites does this exact handle resolve to a real, live profile? There are free tools that check hundreds of sites in seconds — well-known examples are **Sherlock**, **Maigret**, and **WhatsMyName** (tool names change over time, so treat these as examples, not gospel). But remember: the tool only tells you a page exists. It does not tell you the page is your person. You still look with your own eyes.
- **Triage the hits by hand.** Triage means sorting results by what matters. Open each live hit. Is it a real profile with real activity, or an empty placeholder? Some sites show a "taken" page for a name nobody really uses. That is a false hit. Throw it out.

- **Harvest new threads.** On every real profile, collect anything new: a real name, a bio, a linked website, a photo, another handle they mention, a city. Each new item is a fresh thread to pull.
- **Follow the threads.** Feed the new handles back into enumeration. Feed a real name into the next step, permutations. Feed a photo into image search. Round and round, the picture grows.
- **Rank every link on the ladder.** As you go, keep asking: how sure am I that this account is the same person? Write the rung down. Weak, moderate, strong, confirmed. Do not let a weak lead quietly turn into a fact in your head.

Three of those steps deserve opening up, because they are where beginners find the real gold.

Guessing the variations

People do not always reuse a handle exactly. Sometimes they change it a little. `firecat` becomes `fire_cat` becomes `firecat88` becomes `thefirecat`. The trick is that each person changes it in *their own consistent way*. This is their personal naming grammar — their private rulebook for making handles.

Once you spot the rulebook, you can predict handles you have not even found yet.

To generate variations, take the parts you know and mix them:

- First name and last name in different orders: annalee, anna.lee, leeanna, a_lee.
- Add a number they seem to like, or a birth year: annalee92, anna_lee_1992.
- Swap letters for look-alikes, called leetspeak: e to 3, o to 0, so annalee becomes ann4l33.
- Add the small words they favour: theannalee, realannalee, annaleexo.

Then run those guesses back through enumeration to see which ones are real. Predicting a handle feels clever, but a prediction sits on the weak rung until a live profile proves it real.

Turning a username into an email, and back

The part of an email before the @ is called the local-part. In annalee92@gmail.com, the local-part is annalee92. Notice something? That is a handle. So emails and usernames flow into each other.

- **Email to username.** If you find an email, strip off the local-part and run it through enumeration like any other handle. annalee92 might be on ten sites.
- **Username to email.** If you have a handle, guess emails from it. Try annalee92@gmail.com, annalee92@outlook.com, annalee92@proton.me. Then confirm which mailboxes actually exist using an existence check.

An existence check asks a website, in a gentle and public way, "does an account with this email exist here?" Many sites answer that quietly, without ever bothering the account owner. Free tools do this for you across dozens of services. A word of care: a few sites will send the owner a notice if you poke them. Avoid those, or you tip off your subject and cross a line you should not cross. Stay on the public, quiet side of the fence.

How well this pivot works depends entirely on how rare the local-part is. annalee92 pivots beautifully. info or contact or admin tells you nothing, because a million people use them.

The photo is a second thread

A username is thread number one. A profile photo, the avatar, is thread number two. And it is often stronger, because people reuse images even more carelessly than they reuse names.

Two ways to pull this thread:

- **Reverse image search.** You give a search engine the picture instead of words, and it finds other places that exact picture appears. If the same odd photo shows up on three accounts with three different handles, that is a strong sign one person runs all three.
- **Face search.** Some engines find the same *face* across *different* photos. That is even more powerful, because it can link a person even when they used a different picture each time.

But hold the ladder in mind here too.

- **Similar is not identical.** A result that merely *looks* like your photo may be a different crop, a re-post, or a genuine look-alike. Before you trust a match, confirm it is the *same* image — not just a similar one.
- A rare, original photo the person took themselves is a strong signal.
- A stock photo, a celebrity, or a picture thousands of people use is worth nothing for identity.
- An AI-generated face — a fake face made by a computer — will return no matches at all. And that silence is its own signal. A face that exists nowhere else is often a made-up persona.

Do not use face-search tools without written authorisation and a legal check.

Face-search touches people's privacy directly, and in many places it is restricted or outright illegal. Know the rules where you are, get permission in writing, and have a clear lawful reason before you ever run one. This is not optional — it is the one pivot in this chapter that can put *you* on the wrong side of the law.

One case, start to finish

Watch the whole thing work, on a made-up person. Everyone and everything here is invented. Follow the pulls — and notice how little I have to start with.

- **Step 1 — The seed.** A small online shop is getting fake five-star reviews. Every fake review comes from an account with the handle mossyfern. That is all I have. One handle, seven letters, nothing else. I write it down: mossyfern.
- **Step 2 — Enumeration.** I run mossyfern across a few hundred sites. The tool does not hand me a clean list. Here is roughly what it actually prints:

```
[*] Checking: mossyfern
[+] found    photo-site.example/mossyfern
[+] found    gaming-forum.example/u/mossyfern
[+] found    recipe-site.example/mossyfern
[+] found    twitter.example/mossyfern
[-] 404      instagram.example/mossyfern
[!] unknown recipe-site (page loaded, empty?)
4 found, 1 not found, 1 flagged
```

Read that honestly. The tool leans *optimistic* — it counts a page as "found" the moment it loads, even a parked one, so those green lines are hits to check, not leads to trust. Instagram, the one platform I most wanted, has no mossyfern at all. That is not a bug; the handle simply is not there, and I write that down.

Right now this whole run is a weak lead.

- **Step 3 — Triage and harvest.** I open each hit with my own eyes, because the tool only proved a page loads. The recipe site is an empty placeholder — the name is taken, but there are no posts, no avatar, nothing. That is a *false hit*; the tool called it "found" and it is really just a parked, generic page. I throw it out. The Twitter hit is the dangerous one: it is a real, busy

account, but it belongs to a teenage boy in another country posting about football. Same seven letters, completely different human. This is the classic false positive — someone else grabbed the identical handle. I catch it not by the name (the name matched perfectly) but by the *content*, the *photos*, and the *dates*: nothing there touches plants, houses, or a rainy city, and the account is three years newer than the others. Out it goes. Two hits survive. The photo site is real and active: dozens of pictures of houseplants, and a bio that reads "plant mom in the rainy city, she/her." The gaming forum is real too, and the profile links out to a personal blog. Two real profiles. I harvest: a hobby (plants), a hint of location (rainy city), a linked blog.

Same seven letters, different human. That is the homonym trap in the flesh.

- **Step 4 — The avatar pull.** Both real profiles use the same photo: a woman holding a fern by a window. I reverse-image-search it. The exact same photo appears on a small local plant-club page, captioned "our volunteer Deborah Ashworth." Now I have a possible real name. This is a strong signal — a rare, original photo tying the handle to a name. I have climbed two rungs at once.
- **Step 5 — Permutations from the name.** With "Deborah Ashworth" in hand, I build variations: `deborahashworth`, `dashworth`, `deb_ashworth`, `debashworth`. I run them back through enumeration. `deb_ashworth` is live on a professional

networking site, listing a job at a plant nursery in the city of Elmbrook. "Rainy city" and "Elmbrook" and "plants" all fit. Moderate signal on its own, but it stacks neatly on what I already have.

- **Step 6 — The email pivot.** From the blog I found in Step 3, the "contact" page shows `deb.ashworth.ferns@gmail.com`. I do not email it. I run a quiet existence check. That address is registered on the photo site and the gaming forum — the exact two accounts I started with. That is a strong signal, but I stop myself: it is *one* check. One positive result, however clean, earns strong and no higher. It stays strong until a second, independent source agrees. So I hold it there and go looking for that second source.

- **Step 7 — Confirm and stop.** I already have it. The rare original photo tied the handle to the name Deborah Ashworth on the plant-club page (Step 4). The email check now ties the same handle to the same name a second, independent way (Step 6). Two separate sources, arrived at by two separate paths, point to one person. *That* is what lifts me from strong to **confirmed** — not the email alone, but the email plus the photo agreeing. I write up the finding with each rung of evidence laid out plainly, so anyone can check my work. And then I stop. I do not chase her private life beyond the question I was asked. That restraint is part of the craft.

Look back at what happened. Each step handed the next step its fuel. Handle led to profiles. Profiles led to a photo. Photo led to a name. Name led to more handles. Handles led to an email. Email closed the loop. No single clue proved anything. All of them together did.

And notice the mess I did *not* skip past. The tool called a parked recipe page "found," so I opened it and binned it. It matched a stranger on Twitter, so I killed the false positive by reading content and dates, not the name. Real work looks like this.

The traps that catch beginners

Before you go practise, learn the mistakes so you can dodge them.

- **The homonym trap.** Two different people, same handle. This is the most common failure of all. Short and common handles are almost worthless because so many people share them. Trust rare strings, doubt common ones.
- **The recycled name.** Someone abandons a handle; a stranger grabs it later. The old posts and the new posts are two different humans wearing the same name. Verify that the *person* continued, not just the string.
- **The tool that lies.** Enumeration tools break as websites change. They report hits that are not real and miss hits that are. Always look with your own eyes, and run more than one tool, because each checks a different list of sites.

- The fake face. A stock photo or an AI face links nothing to a real identity. Do not build a case on a picture anyone could use.

Every one of these traps has the same cure: climb the ladder and gather independent clues.

Exercise

Do this on yourself. It is the safest and most eye-opening way to learn.

- Write down the oldest username you can remember using — the one from your teenage years, the one you typed everywhere without thinking.
- Run it through a free cross-platform username checker. (A note on tools: some are just websites you open; others are small free programs you install, sometimes run from a terminal. If a terminal is new to you, start with a browser-based checker — the method is identical, and you never need to pay to practise anything in this book.) See how many places it still lives.
- Open each real hit. Ask: what could a stranger learn about me from this alone? A hobby? A city? A real name? An old photo?
- Now try three variations of that handle, using your own naming grammar. How many are also you?

- Finally, take the oldest profile photo you can find of yourself and reverse-image-search it. Where else does it appear?

Sit with what you find. This is exactly what someone would do to a target. Feeling it from the inside teaches you both how the method works and why we treat it with such care.

Turn it around. The fix for everything you just found on yourself: stop reusing one handle everywhere, retire old accounts you no longer touch, and keep your real name off throwaway profiles. That single reused username is the thread that ties your whole life together — cut it. (Full checklist: Appendix D, "Take it back.")

Recap

- A username is often the single best starting thread, because people reuse the same handle across many sites for years, and one handle leads to another, then to photos, names, and emails.
- Work the thread in order: enumerate where the handle exists, triage the real hits by eye, harvest new clues, guess variations, and pivot between username and email — with the profile photo as a strong second thread.
- Matching names never prove same person. Climb the confidence ladder from weak to confirmed, stack independent clues, watch for homonyms and recycled names,

and stay strictly legal and open-source.

CHAPTER SIX

Email

The master key

Your own email address — count how many places you typed it this year. A shopping site. A newsletter. A new app. A work login. A game your kid wanted. A dentist's booking form.

Now consider how long you have had that address. Many people keep the same email for ten years, some for twenty. We change our phone number when we switch carriers. We move house. We change jobs. But the email stays.

That is why, in OSINT, the email address is the master key.

Quick reminder of the word. OSINT means *Open-Source Intelligence*. It is the craft of finding facts about a person or thing using only information that is already public — the open web, public records, things people posted themselves. No hacking. No tricks that break the law. Just careful looking.

An email address is one small string of text. But it is a *stable*, *unique* label for one person. *Stable* means it rarely changes. *Unique* means only one person uses it. Put those two together and you have something rare: a single dot that connects to dozens of other dots. Where the person shops. What apps they use. What their real name is. Sometimes their photo. Sometimes

an old password habit they never fixed.

In this chapter I will show you how to pull on that string — gently, legally, and step by step.

An email address is the one label a person carries for years across the whole internet. Find it, and doors start to open.

Why an email ties so much together

Every account you make asks for an email. It has to. That is how the site says "you are you" when you log back in. That is how it sends you a receipt or a reset link.

So your email quietly becomes a list of everywhere you have been. Not on purpose. It just happens.

For an investigator, that means one email can answer questions like:

- Which websites and apps is this person signed up on?
- What is their real name?
- What username do they like to use?
- Has their data ever leaked in a hack — and if so, from which service?
- Is this the *same* person as that other account I found?

We will take these one at a time. But first, the rule that keeps you safe and honest.

The line you never cross

There is a bright, clear line in email OSINT. Learn it now and it will protect you your whole career.

You may check whether an account exists. You may never touch the account.

Checking is fine. Asking a website "is there an account here for this email?" is a public question anyone can ask. Reading what the website *chooses* to show you is fine.

Touching is not fine. Logging in is not fine. Guessing a password is not fine. Using a leaked password to get inside an account — even to "just look" — is not fine. That is not OSINT. It is breaking in, and in most countries it is a crime.

The simplest way to hold it in your head: picture standing on the pavement outside a house. You are allowed to notice the lights are on, the car is in the drive, the mailbox says "Karas." You are *not* allowed to open the door. Same street, different act.

Confirm that an account exists. Read what the site volunteers. Then stop. Never log in, never guess, never use a leaked password.

We will come back to this line more than once, because it matters most where the temptation is strongest.

Method 1 — Where is this email used?

The first thing you want from an email is a map. Where is this person signed up?

There is a clever, gentle way to find out, and it uses a feature every website already has: the "forgot password" and "sign up" forms.

The idea, in plain words: when you go to a website and start the "forgot password" flow, the site has to ask itself: *do we have an account for this email?* And then it usually tells you the answer, right there on the screen, before you have proved anything.

- If it says "we've sent a reset link to your email" or "that email is already in use" — an account exists.
- If it says "we couldn't find an account with that email" — no account.

That difference in wording is called a **differential**. It is just a fancy word for "the site gives two different answers, and the answer tells you what you want to know."

Notice what you did *not* do. You did not log in. You did not get the reset link (that goes to the target's inbox, not yours). You did not touch anything. You asked a public question and read the answer. That is why this technique is on the safe side of the line. But be honest about what "passive" means: it means you did not log in — not that you are invisible. A site can rate-limit these checks, log the attempt, or count it as a breach of its terms of service. Quiet is not the same as undetectable.

You could do this by hand, one website at a time — click "forgot password," type the email, read the answer, write it down. Doing it a few times by hand teaches you what the answers look like.

There are also tools that check many sites at once and hand you a list. Their names come and go, so I will not marry you to any one of them. A well-known free one has been *Holebe*, which checks over a hundred sites and reports whether the email is registered on each; a hosted service called *Epieos* does something similar in your browser. Treat these as *examples of the method*, not as gospel. Email tools rot fast — they depend on other websites' forms, and those forms change every month. Always cross-check a surprising "not found" with a second tool before you believe it.

What you get at the end is gold: a footprint. "This email is on GitHub, Spotify, and an old forum." Each is a new thread to pull.

Method 2 — Turning a name into an email (the pattern trick)

Sometimes you have the opposite problem. You know a person's name and where they work, but you do not have their email. You want to *work it out*.

One key fact makes this possible. Companies almost always follow one pattern for everyone. They do not invent a new style for each employee — that would be chaos for their IT

team. So they pick one recipe and stamp it on the whole company.

Common recipes look like this. Say the company domain is brightmail.com and the employee is Jason Karas:

- first.last → jason.karas@brightmail.com
- flast (first initial + last name) → jkaras@brightmail.com
- first → jason@brightmail.com
- f.last → j.karas@brightmail.com
- firstl → jasonk@brightmail.com

And here is the beautiful part, worth saying out loud:

Find one confirmed email at a company, and you have cracked the pattern for everyone there.

If you can confirm — from a public source — that maria.velez@brightmail.com is real, you now know the recipe is first.last. So the boss you are looking for, Jason Karas, is a **strong** prediction: jason.karas@brightmail.com.

Where do you find that one confirmed address? Public places, all open:

- A company press release or a PDF report (try a search like "@brightmail.com" filetype:pdf).
- A staff "contact us" page.

- A public code repository, where developers' emails sit in the commit history.
- Domain-lookup tools that already catalogue a company's email format.

A small warning, so you look like a professional and not a spammer. Once you have your predicted address, **do not email it to "test" it**. Sending mail tells the person someone is poking around, which defeats the whole quiet purpose. There are ways to check if an address is real *without* sending anything, and tools built exactly for that — but for a beginner, the honest move is: find one confirmed address, apply the pattern, and treat the result as **strong** — not **confirmed**.

One more note. This pattern trick is cleanest for *workplaces*. For personal Gmail-style addresses, people mix in nicknames, birth years, and random numbers, so the guesses explode. Keep it for company domains, where discipline rules.

Method 3 — Has this email been in a breach? (exposure only)

Big companies get hacked. It is sad but ordinary. When they do, lists of their users' emails — sometimes with passwords, dates of birth, or addresses attached — leak onto the internet. We call each event a breach.

For an investigator, breach *exposure* is a useful, careful signal. Note that word: *exposure*. We are checking *whether* an email

showed up in a known leak, and *which* leaks. We are not going hunting for the stolen passwords, and we are certainly not using them.

The gold-standard free tool here is Have I Been Pwned (often written HIBP). You type an email, and it gives a simple list: this address appeared in the LinkedIn breach of such-a-year, the Dropbox breach of another, and so on. It does *not* show you the passwords. It shows *presence* — the fact of exposure, and the names and dates of the leaks.

Why is that useful to an OSINT investigator? Two reasons, both clean:

- It confirms the person had an account somewhere. If the email shows up in the "MyFitnessPal" breach, then this person used MyFitnessPal — one more site on your footprint map, sometimes for a service that has since shut down.
- It hints at timing. The breach has a date, which tells you the account existed *by then*. It sketches a rough timeline of the person's online life.

The line applies here too. Checking *exposure* — "yes, in these breaches" — is public knowledge about a public event, so it is safe. Reading, or using, the leaked passwords behind it is the door you already know you never open.

Breach data is also *messy*. Leaks get mislabelled. Old data gets recycled and re-sold under new names. So treat any breach hit as

a lead, not a fact. It points you somewhere. It does not prove anything on its own.

And here is the part beginners get wrong, so learn it once and carry it everywhere: a zero result does not mean the address is safe, or fake, or unused. It means the tool has not indexed a leak that holds it — nothing more. Plenty of real, heavily-used addresses come back empty: their breaches were never added, or the address is newer than the leaks. "Zero breaches" reads as *not found here*, never as *clean* — the absence of a positive, not the presence of a negative. Treat it as a small nudge, never as a verdict.

What these checks actually look like on the screen — one hit, one empty:

```
t.orsini@northvale-labs.com
Breaches: 1 (PixelForge Design, 2019)
Pastes: 0
---
info@northvale-labs.com
Breaches: 0 Pastes: 0
"0" != "safe" -- may just be unindexed
```

A breach check tells you where and roughly when someone had accounts. It never gives you permission to open one.

Method 4 — The Gravatar pivot

One lesser-known trick feels like magic the first time it works.

There is a service called **Gravatar** — short for "globally recognized avatar." An *avatar* is just the little profile picture next to your name. Lots of websites (comment sections, developer sites, blogs) use Gravatar so that one photo follows you everywhere. You set it once, and it shows up across all of them.

And here is the quirk we can use: Gravatar links your photo — and sometimes a whole little profile — to your email. So if you have someone's email, you can sometimes pull back their photo, chosen username, real name, location, and links to their other accounts. From nothing but an email string.

You do not need to understand the math to use it. In simple terms: you take the lowercased email, run it through a standard fingerprint (called a **hash** — a fixed-length code made from the text), and ask Gravatar's public address for the matching picture and profile. Many tools do this step for you, so you just paste the email and read the result.

What you get, when it works: a face to put to the name, a username to feed into your other searches, and sometimes a location. Not everyone has a Gravatar, and many set only a photo. But when it hits, it hits big — a photo and a username from a single email is a wonderful thread to pull.

Method 5 — Pivoting: email to username to name and back

Everything above shares one shape. You start with one piece of information, and it hands you another, which hands you another. This is called **pivoting** — swinging from one fact to the next, like crossing a river stone by stone.

The three stones you will hop between most often are **email**, **username**, and **name**. Here is how they feed each other:

- **Email** → **username**. The breach check or the Gravatar profile hands you a username the person likes to use, say `jkaras_dev`.
- **Username** → **name**. You search that username across social sites and forums (there are whole tools for this, covered in another chapter) and find a profile that shows a real name: Jason Karas.
- **Name** → **email**. Now you apply the company pattern trick from Method 2 and predict `jason.karas@brightmail.com` — and confirm it exists with the gentle "forgot password" check from Method 1.

See the circle? Each fact tightens the picture around the next. This is the real heartbeat of OSINT. No single search "solves" a person — you build them up, one confirmed stone at a time, checking each new fact against another source before you trust it.

A worked case: one email opens a picture

Watch the whole thing as one small, *made-up* story. Every name and address here is invented — follow the moves, not the details.

You are handed a single email address:

t.orsini@northvale-labs.com. That is all. Your job is to work out who this is.

- **Step 1 — Read the address itself.** Before any tool, just look. The domain is northvale-labs.com, a company. The local part is t.orsini, which smells like a f.last pattern — first initial t, last name orsini. So you are looking for someone named T-something Orsini at Northvale Labs — a weak read for now, but a free one, and you got it just by looking.
- **Step 2 — Map the footprint.** You run the email through a gentle "forgot password" check across common sites, using the method from earlier. A code-hosting site says "we've sent a reset link" — a clean yes, exists. But the professional network gives you *the same polite message no matter what you type*: "if an account exists, we'll email a link." That is deliberate. Many big sites now hide the differential exactly so people can't do what you're doing. So you do **not** get to conclude "she's on this network." You mark it *unknown*, not *yes* — no over-reading a message the site engineered to tell you nothing.
- **Step 3 — Check the exposure.** You look the email up on Have I Been Pwned. It appears in one breach, from a design-tools service, dated a few years back (the PixelForge Design, 2019 line in the box above). That confirms an old

account elsewhere and roughly dates their online life. You write it down as a *lead* — nothing more. You do not go looking for any password. Out of habit you also check the company's generic `info@northvale-labs.com` — and it comes back zero breaches, zero pastes. You note it as a non-result, not a clean bill — again: a blank is not "clean."

- **Step 4 — Try the Gravatar pivot.** You feed the email to a Gravatar lookup. A profile comes back: a photo, a username `torsini`, and a listed city. Now you have a face and a handle from one address.
- **Step 5 — Pivot the username to a name.** You put the email in quotes and search it: "`t.orsini@northvale-labs.com`". The first page is mostly noise — those grey "email lookup" and "find anyone" sites that scrape an address, wrap it in ads, and know nothing you don't. You skip every one; they are junk dressed as results. Below the noise sits one real hit: a conference attendee list (a public PDF) naming *Teresa Orsini*, Northvale Labs. Then you pivot the handle instead — `torsini` across social and forum sites — and a second profile confirms the same name and employer. Two independent sources, so now you believe it. That matches your Step 1 guess. The stones connect.
- **Step 6 — Confirm the pattern.** Just to be sure of the recipe, you find one *other* Northvale Labs address in a public PDF on their site: `m.dubois@northvale-labs.com`. Same f.last shape. So the company pattern is confirmed, and

t.orsini@northvale-labs.com sits neatly inside it — your original email verified from a second direction.

- A snag that does not resolve. On that same page you also spotted orders@northvale-labs.com and orsini.family@example-mail.com. Neither belongs to *one* person. The first is a role mailbox — a shared inbox a whole team reads. The second is a household account, shared by a couple or a family. Run it through Gravatar and get back one photo — whose face is it? You cannot tell. It could be Teresa, a partner, or an adult child who set up the account years ago. The address maps to a *household*, not a *human*, and no amount of pivoting splits it back apart. So you write down what is honestly true and no more: *shared address — could be the owner or a family member — cannot separate them from here*. That is where this thread ends. Not a clean answer. A partial one, marked as partial.

From one string of text, you now have a name, a face, a city, an employer, a rough timeline, and two confirmed accounts — plus one channel marked *unknown* and two addresses you refused to pin on a single person. Every step public, every step gentle, never once touching an account. Some of your sharpest moves were the ones where you *didn't* conclude: the hidden differential, the empty breach, the shared mailbox.

A negative or ambiguous result is real data. "Unknown" written down honestly beats "yes" guessed wrongly, every

time.

Exercise: check your own exposure

Do this on **your own email** — the one that belongs to you. This is the safest possible practice, because you are the subject and you have every right.

1. Open Have I Been Pwned in your browser. 2. Type your own email address and run the check. 3. Read the list of breaches it comes back with. For each one, notice two things: *which service* leaked, and *what year*. 4. Write a short timeline for yourself: "I had an account on X by year Y." Notice how much of your own online history this sketches out — from one address. 5. As a bonus, look at *what kind* of data each breach exposed (some list "passwords," some "dates of birth"). If any breach exposed a password you still use anywhere, go change it. You have just turned an OSINT lesson into real self-defence.

Sit with how much your single email revealed. That feeling — "wow, it's all connected to this one thing" — is exactly what you now understand about everyone else's email. Use it with care.

Recap

- An email address is the master key because it is *stable* and *unique* — one person keeps it for years, so it quietly links to everywhere they have signed up, their real name, their photo, and their online history.

- The core moves are gentle and public: map where an email is used with the "forgot password" reveal, work out a company's email pattern from one confirmed address, check breach *exposure* on a service like Have I Been Pwned, try the Gravatar photo pivot, and swing between email, username, and name.
- Never cross the line you set at the start: confirm an account exists and read what a site volunteers, then stop. Logging in, guessing, or using a leaked password is not OSINT — it is a crime.
- Facts feed each other: email, username, and name form a loop, and each new fact is only worth trusting once a second source confirms it. A role mailbox or a household address maps to a group, not a person — write down what is honestly true and no more. "Unknown," written down honestly, beats "yes" reached wrongly.

Turn it around. Shut this pivot down on yourself: use a unique address (or an alias) per site, turn on two-factor, and don't reuse one email as your username everywhere. See Appendix D (Take it back).

CHAPTER SEVEN

Phone Numbers

The personal anchor

Your own phone number, for a second.

You have had it for years. Maybe you have had it for a decade. You gave it to your bank. You gave it to your doctor. You put it on a job application. You typed it into a food-delivery app at 2 a.m. You saved it in a hundred friends' phones under your name.

A phone number is one of the most personal and most stable clues a person leaves behind. People change their email. People delete old social accounts. But most people keep the same phone number for a very long time. That is why I call it *the personal anchor*. It holds a person in place.

This chapter is about how to read a phone number, and how to safely see what is publicly tied to it. We will move slowly, and every term gets explained the moment it appears. Stay calm and careful the whole way: a phone number is not just data. It belongs to a real person.

What a phone number actually is

One confusing thing to clear up first.

There is no single official "reverse phone book" for the world. You cannot type a number into one magic website and get the owner's name, verified and true. Anyone who tells you otherwise is selling something, and often selling stolen data.

What we actually do is smaller and honest. We read the *structure* of the number. We check where it *shows up* in public. And we check which *apps* the number is registered on. Each of those is a small clue. We stack the clues. That is the whole game.

A phone number is a lead, never a conclusion. One bit points somewhere. It does not prove anything by itself.

Reading the number itself

Before you look anything up, you read the number. This part is completely safe. You are not touching any network. You are not pinging anyone. You are just looking at the digits like a mapmaker looks at coordinates.

First, put the number into a clean international form. Investigators call this E.164 — that is just the name of the standard format. It looks like this:

+ then the country code then the rest of the number, with no spaces or dashes.

A made-up United States number written cleanly is +1 202 555 0134. A made-up United Kingdom number is +44 20 7946 0958. The + means "international." The first digits after it tell

you the country.

- **Country code** — the leading digits tell you the nation. +1 is the US and Canada. +44 is the UK. +380 is Ukraine. +49 is Germany. This is the single most useful thing you learn, because it tells you *which country's rules apply* to everything you do next.
- **Area or region prefix** — the next block often points to a city or region, especially for landlines (a landline is a fixed home or office phone, not a mobile). For mobile phones this block points to the carrier that *first* issued the number.
- **Line type** — many free tools will also guess whether it is a mobile, a landline, or a VoIP number. VoIP means "voice over internet" — a phone number that lives on the internet, like Google Voice or a free texting app. This matters a lot. A VoIP number is cheap and anonymous. It belongs to no fixed person or place. If a number that claims to be someone's personal mobile turns out to be VoIP, be extra careful.

There is one trap here I must warn you about right away.

The carrier a tool shows you is usually the **original carrier** — the company that first handed out the number. But people can *keep their number and switch companies*. This is called **number portability**. So the number may say "Carrier X" while the person actually moved to Carrier Y years ago. Never write down "this person uses Carrier X" from a static reading alone. It is a weak reading, not a fact. This is the one place in the book that owns

portability — when it comes up again later, I will just point back here.

Reading the number offline tells you country, region, and line type. It never tells you the current owner.

The tools for this part come and go, and their names will change after this book is printed — most wrap the same open-source engine that validates numbers. Do not memorise any tool. Remember the *steps*: clean it to +country code number, then read country, region, and line type.

Seeing what the number is publicly tied to

Now we go looking. Safely.

Search the number in quotes across the web

This is the simplest move, and it is often the fastest path from a bare number to a real name or a business. People post their number all over the open web — in classified ads, business listings, forum signatures, old CVs, marketplace posts.

Type the number into a search engine inside **quotation marks**. Quotes tell the engine "find this exact thing." So you search "555 0134".

Here is the important part. People write their number a dozen different ways. So you search it a few different ways too:

- With the country code: "+1 202 555 0134"

- In the local style: "(555) 0134"
- With dashes: "555-0134"
- With spaces or dots

One format will hit a page that the others miss. When you find a page that lists the number, it almost always lists something *next* to it — a name, an email, a shop, a username. That neighbour is your next step.

Try more than one search engine. Google is not the whole internet. Bing and Yandex index different regional pages. For numbers from Eastern Europe and the region around it, Yandex is often stronger.

The trap here: **number reuse**. Phone numbers get recycled. When a person drops a number, the carrier gives it to someone new later. So an old page might show the *previous* owner. And businesses often share one number across many staff. Never assume the page and the person are a perfect match. Look for a second source that agrees.

Crowdsourced "who called me" name apps

You have probably seen apps that claim to tell you who is calling. Truecaller, GetContact, and similar apps are the well-known examples. Here is exactly how they work, because understanding this protects you from a big mistake.

These apps do not have a secret official directory. They show you how other people saved this number in their own phones. When many users let an app read their contact list, the app learns that lots of people saved +1 202 555 0134 as, say, "Mike Plumber." So when you look it up, it shows "Mike Plumber."

These apps cut both ways. The moment you install one and let it read your contacts, you hand over everyone in your phone — and your own number becomes searchable to others.

Sometimes this is genuinely helpful. One app, GetContact, is interesting because it shows you *several* names that different people saved for the same number. You might see "Mike Plumber," "Mike Landlord," and "SPAM." That cluster of names is a little story. It tells you the person's roles and how others see them, not just a single label.

But now the warning, and this is the most important warning in the whole chapter.

Crowdsourced caller-name apps can lie, and they lie constantly. A name in one of these apps is a rumour, not a fact.

Why can they lie? Because the names are typed by random strangers. Someone can save your number under a joke name. Someone can save it wrong. An old contact can be years out of

date. The number may have changed owners. The app just repeats whatever people typed. Bellingcat, a respected investigative group, warns plainly that these names have "mixed accuracy" and depend on how the number appears in other people's phone books.

So treat every returned name as a *separate small clue*. A name means much more if it shows up independently in two different apps or sources. A single name from a single app proves nothing.

One more thing, for your own safety. Many of these apps quietly upload *your* contact list when you use them. If you look someone up, you may hand over all your friends' numbers, and you may become searchable yourself.

The strongest beginner move: number to messenger to username

Now for the single most productive path in phone work — the move that turns a lonely number into a real trail.

The core idea, in plain words.

A phone number is often the *key* to a messenger account. Messenger apps like WhatsApp, Telegram, Signal, and Viber let you find people by their phone number — that is a normal, built-in feature so friends can find each other. When you check a number in one of these apps, you may see a **profile**: a photo, a

display name, a short bio, and sometimes a username.

That username is the treasure.

Because once you have a username, you are back in territory you already know. In an earlier chapter we learned that people reuse the same username across many sites. So a username you pulled from a messenger becomes a search you can run everywhere else. The chain looks like this:

number → messenger profile → username → the rest of that person's online life.

Of all the paths in this chapter, Telegram is often the richest, because it can show a real @username even when it hides the actual phone number. And an @username is exactly the kind of clue that pivots across the whole web.

A gentle but firm caution. Checking a number in a messenger usually means saving it as a contact, and saving a contact can notify the person — they may see a new contact appear, or you may pop up in their "contacts who joined" list. So, as Chapter 3 covered, do this from a clean research account, never your own. And remember privacy settings: if someone set their app to "nobody can find me by my number," the check simply comes up empty. That is normal.

A number often unlocks a messenger username. A username unlocks everything you already know how to

search.

A worked case: from a number to a name

Here is one fictional example, start to finish, so you see the clues stack. Every name and number here is invented.

You are trying to identify who posted a suspicious "car for sale" ad. The only contact detail in the ad is a phone number: +1 202 555 0177.

- Step 1 — Read the number offline. You clean it to +1 202 555 0177. The +1 tells you it is a US or Canada number. A free offline parser gives you this:

```
+1 202 555 0177
valid:      yes
line type:  mobile
region:     Washington DC (origin only)
carrier:    "Verizon" (original issuer)
```

An offline parser returns only these four things — country, region, line type, original carrier — never a name. The region is only where the number was *born*, not where the person lives now. And the carrier caveat applies: it is the *original* issuer, so you do not write "uses Verizon" anywhere. Write down only what is safe: "country = US, mobile, region = origin not home, carrier = unreliable (portability)." Now you know which country's web and search engines to lean on.

- **Step 2 — Search the number in quotes.** You search "202-555-0177", then "+1 202 555 0177", then "(202) 555-0177" on more than one engine. The dashed version hits an old post on a hobby forum where a user signed off with that number and the handle *desertdrew*. Interesting. But it is one page, so it is only a weak lead. (Had the parser tagged this VoIP — a Google Voice or free-app number — it would place *nobody* anywhere: no home region, no real carrier. You would grade the whole location angle weak and not build on it. A dead signal is fine; pretending it is strong is not.)
- **Step 3 — Check crowdsourced name apps, carefully.** Here is where the first app tries to fool you. It returns one loud name:

```
+1 202 555 0177
saved by others as: "Mike Plumber"
(crowd label from strangers' address books)
```

If you were sloppy, you would write "the seller is Mike, a plumber" and go home wrong. But stop. A single crowd name is weak on its own — a rumour typed by strangers, not a fact.

"Mike Plumber" does not match the forum handle *desertdrew* at all. Most likely the number was recycled and "Mike Plumber" is the previous owner, still living in strangers' address books years after he dropped it — the classic wrong answer a naive reader accepts and stops on. So you keep "Mike" as a weak, probably-dead lead. Then you check a *second* app: it returns "Drew" and "Drew Marketplace." A third app also returns

"Drew." Now *two independent apps* agree on "Drew," and it lines up with the forum handle *desertdrew*. Corroboration lifts "Drew" from *weak* to *moderate* — but only because you refused to trust the first single name.

- **Step 4 — Check messengers.** You save the number as a contact and open Telegram. It resolves to an account showing a photo, the display name "Drew," and the username *@desertdrew*. Careful here — the owner *set* that name and username themselves, so they are *not* an independent source. The same person controls both the number and the name; a self-chosen name only corroborates that this account uses the handle *desertdrew*, it cannot confirm who the person is. So this does not lift "Drew." It stays at *moderate*, held there by the two independent apps and the forum post — sources the owner did not write.
- **Step 5 — Pivot the username.** You take *@desertdrew* and search it across other platforms, exactly as we practised in the username chapter. It appears on a photo-sharing site and a review site, both under the same handle, both placing him in the Washington DC area from Step 1. *These* are independent — the owner did not control what a review site or someone else's page says — and their agreement is what finally lifts "Drew" to *strong*. You now have a coherent picture: one person, one handle, one region, reachable at one number.

Notice what you did *not* do. No call, no login trick, no stolen data. You read public structure, searched public pages, used a

normal app feature, and pivoted a public handle — confirming each lead against an independent source before trusting it.

One advanced trick, and why beginners should skip it

You may eventually read about a clever technique where an investigator starts a "forgot my password" flow on a website, sees a hint like "we sent a code to the number ending in 77," and stitches together many such hints from many sites to rebuild a full hidden number.

I am naming it so you recognise it. I am also telling you to leave it alone for now.

Poking at a live login or account-recovery system to squeeze out account information sits in a legal grey area. Depending on your country and your intent, it can cross computer-misuse laws and app rules. It can also send a security alert to the real owner. This is not beginner ground. As a foundation-level investigator, you get remarkably far with the safe steps above and never need to touch it. If your work ever genuinely requires it, that is a moment for training, written authorisation, and a careful legal check — not curiosity.

If a technique makes you interact with someone's login or recovery system, treat it as sensitive and step back.

Foundations first.

Exercise: safely look up your own number

The safest number in the world to practise on is your own. You already know the truth about it, so you can *grade* what the tools tell you. That makes it the perfect teacher.

Do this:

- **Write your number in E.164.** Put it in the + country code number form, no spaces or dashes. Confirm you got the country code right.
- **Read it offline.** Use any free number-reader to check country, region, and line type. Does it match reality? If it lists a carrier, notice whether it is your *current* carrier or an old one — this is how you feel portability with your own eyes.
- **Search it in quotes.** Search your own number in a few different formats, in more than one search engine. Where does it show up? You may be surprised what is public. Anything you find, someone investigating you could find too.
- **Check a caller-name app.** Look yourself up in one crowdsourced app. What name comes back? Is it right, wrong, funny, or blank? Sit with how *unreliable* it is even for a number you know perfectly.
- **Write two or three sentences** on what was accurate, what was wrong, and what surprised you.

That reflection is the real lesson. Feeling the tools be wrong about *you* is the fastest way to learn to never trust them blindly about a stranger.

Recap

- A phone number is a personal, long-lasting anchor. You read its structure offline (country, region, line type), you search it in quotes in several formats across several engines, and you check which apps it is registered on — every hit is a lead to confirm, never a fact.
- Crowdsourced "who called me" name apps show how strangers saved a number, so they are easily wrong and easily faked. A name only counts when two independent sources agree.
- The strongest beginner move is number → messenger profile → username, because a username drops you back into cross-platform search you already know. Do it from a clean research account, and leave login-and-recovery tricks alone as a legal grey area.

Turn it around

Everything in this chapter can be pointed at *you*. Turn off "find me by my number" in your social apps and messengers, delete any caller-name apps — they upload your contacts and make you searchable — and consider a separate number for signups. Appendix D (Take it back) walks through it.

CHAPTER EIGHT

Your First Real Investigation

Putting it all together

This is the last chapter of the book. So let us do the thing the whole book was for.

We are going to run one real investigation, start to finish. One clue in. One honest little report out. Everything you learned in the earlier chapters gets used here — the search skills, the username hunt, the email checks, the phone checks, the confidence ladder, the OPSEC habits, and the ethical line you do not cross.

I invented this whole case. Every name, username, email, and phone number below is made up for teaching. None of them are real people. But the *method* is exactly how a professional works. Follow along slowly. Do not rush to the answer. In this work, how you get there matters as much as where you land.

One thing before we start. Keep your notes open the entire time. A blank text file is fine. Write down what you searched, what you found, when you found it, and *how* you found it. If you cannot trace a conclusion back to a note, you cannot defend it later. So we take notes from the very first click.

And one habit that saves cases: capture the page the moment you see it. Web pages change and get deleted — often the interesting ones, often right after you find them. So when something matters, take a screenshot, copy the exact web address, and save the page to a public web archive so there is a frozen, dated copy you can point to later. A note that links to a live page is only as reliable as that page tomorrow. A note that links to an archived copy is evidence.

A finding you cannot trace back to a note is not a finding. It is a guess.

The Setup

Here is our clue. Just one.

A small online marketplace for handmade goods got a complaint. A seller there, using the shop name "pinewicketcrafts", took money for an order and vanished. No delivery. No replies. The buyer asks a simple, fair question: *is there a real person behind this shop, and is this shop connected to other shops that might be doing the same thing?*

That is our objective. Notice how narrow it is. We are not "researching a person." We have a clear, answerable question. That focus is the whole job. A loose question makes you wander forever. A tight question tells you when to stop.

Write the objective at the top of your notes, word for word:

Objective: Is there a real, findable person behind the shop "pinewicketcrafts," and is that shop linked to any other seller accounts?

Everything we do now serves that sentence.

The Investigation

- Step 1 — Set up clean, then look before touching. Before I search anything, I open my research browser, not my personal one, signed out of my own accounts. No logging in, no messaging, no "contact seller" — passive work, public pages only. Then I write my first note: date, time, objective, and the single seed I have — the shop name pinewicketcrafts.
- Step 2 — Start with the plainest search there is. I put the shop name into a search engine, in quotes, exactly: "pinewicketcrafts". Quotes force an exact match, so I do not get pages about pine trees and cricket. Two things come back. First, the marketplace shop page itself. Second, a tiny result from a crafts forum where someone posted with the same name. Both go into my notes with their full web addresses and the time I saw them. I do not trust either yet. Right now they are *leads*, nothing more.
- Step 3 — Read the shop page like a document, not a story. The shop page is thin, but thin pages still leak. In the "About" section the seller wrote one friendly line: *"Made with care in the workshop by Jordan."* A first name. Good, but weak

— a first name alone points at millions of people. I also note the shop opened eight months ago and lists a general region, not an address. I write down every fact and, just as important, I write down what is *missing*: no full name, no phone, no email on the page. Negative results are real results. They show how far I looked.

- Step 4 — Pivot on the username. Here is the first real pivot. People reuse handles. The same name they picked on one site, they often picked on ten others, because it is easy to remember. So I take pinewicketcrafts and check it across many platforms. (In an earlier chapter you saw the free username-search tools that do this — they simply ask each site, "does this name exist here?") The result is not a tidy list. It never is. Some sites say "not found," one rate-limits me, one loads a blank page, and two time out — the tool flat-out warns me which ones it could not read:

```
enum v1.4 · 312 sites · 14:22
[*] pinewicketcrafts
[+] found    marketplace.example/shop/pinewicketcrafts
[+] found    woodcraft-forum.example/u/pinewicketcrafts
[+] found    photos.example/pinewicketcrafts
[-] 404      github.example/pinewicketcrafts
[!] 429      reddit (rate-limited, retry later)
[?] 200?     blogs.example (page loads but blank)
[!] timeout 2 sites unreachable
3 found · 1 not found · 4 unclear
```

So I do not have "the accounts." I have three *candidates* and a blind spot — a rate-limited site, a blank page that could be

anything, and two that timed out. Four sites I could not actually read. I write that blind spot down, because a gap I know about is safer than a clean-looking list I trust too much. Same name is a *lead*, not proof. That caution is the difference between an investigator and a rumour.

- Step 5 — Confirm each candidate, and throw one out.

Now I open the three candidates and look, instead of counting names. Two line up fast: the woodworking-forum profile talks about birdhouses, the marketplace shop *sells* birdhouses, same photo style, overlapping dates. Good — that is two independent places agreeing. But the third, the photo account, is wrong. I almost took it — the handle matched exactly. Then I actually read it: it is a teenager posting gym selfies and sneaker pics, last active years ago, nothing about wood at all. Someone else grabbed the same word first. *Same handle, different human*. I strike it from my notes and write why: *"photos.example/pinewicketcrafts — REJECTED, unrelated person, no craft content, matched on name only."* That rejection is not wasted work. It is the work. On the confidence ladder, two independent matches that survive a real look beat three that I never checked. I now put it at the moderate rung — the marketplace shop and the woodworking profile are probably one person — and I keep the door open, because I am still missing those two unreachable sites.

• **Step 6 — Pivot from a username to an email.** On the hobby-woodworking profile, in a public post asking for advice, the person signed off with a contact address so buyers could reach them: `jordan.pinewicket@example-mail.com`. That is a gift, and it is public, so I am allowed to note it. A new identifier means a new door. I run the safe, standard email checks you learned earlier — not to break into anything, but to see what is *publicly* tied to it. I check whether the address shows up in known public breach lists (a legitimate "has this email been in a leak" service), and I do a plain quoted search for the address itself. Here is what the breach check hands back — and notice the second line, because I also ran the shop's generic contact address to compare:

```
jordan.pinewicket@example-mail.com
Breaches: 2    (Collection-2019, ForumX-2021)
Pastes:  0
---
info@pinewicketcrafts.example
Breaches: 0    Pastes: 0    (no records)
note: "0" is not "clean" — may be unindexed
(2nd lookup delayed — rate-limited 6s)
```

Two things to read carefully. The personal address has been around for years and turns up in old breaches — so it is a genuine, used mailbox, not a throwaway made yesterday. The *shop* address returns nothing at all — but, per Chapter 6, a blank is not clean. I note it and move on. The quoted search then turns up one more real thing: the personal email left a public

review on a *second* marketplace shop years ago, a shop called "oakwicketgoods."

- Step 7 — Follow the new lead, carefully.

"oakwicketgoods." Different word, same odd middle piece — *wicket*. That echo is interesting, but an echo is not evidence. I do not declare the shops connected because two names rhyme. I go and look. The second shop is closed now, but its old public page still lists the same region and the same birdhouse photos. Same email in the review. Same craft. Same style. That is now two shops tied together by an email address and matching content — real corroboration, not a hunch. I write it down plainly: *two seller accounts linked by one public email and matching product photos; moderate — two independent ties so far, with nothing against them.*

- Step 8 — Pivot toward a phone, and read it honestly. The old "oakwicketgoods" page had one more scrap: a public post from years back where the seller offered a phone number for local pickup — +1 202 555 0147. I run the *passive* phone checks from the earlier chapter. And the lookup is a good reminder that a number lies as often as it talks:

```
+1 202 555 0147
valid:      yes (format only, not assignment)
line type:  VoIP (not a home carrier)
region:     Washington DC (number's origin)
carrier:     virtual / app-based
```


Read that and cool down. A VoIP number is an app, not a doorstep — its "region" is where the number was *born*, not where its owner sleeps. So this does not place Jordan anywhere. It is a weak signal, and I grade it weak. What it *is* good for is one thing: a plain search of the number turns up the same digits on both shop pages — so it is a *third* independent thread between the two shops, on top of the email and the photos. That third thread is what lifts the shop-to-shop link from moderate to **strong**; the number itself still stays weak as a way to place a person, and that is all I let it do. Then the same search quietly surfaces something else — a personal social profile, in the name "Jordan Alvarez," with home-life photos, a partner tagged, and what looks like a house and street in the background. I note that surname, but I do not lean on it: it hangs off this one weak VoIP number and nothing else has confirmed it. A single weak lead does not earn a place in my conclusion — so "Alvarez" stays a flagged, uncorroborated maybe, and my finding stays with what I can stand behind: a seller who goes by **Jordan**. And here I stop.

- **Step 9 — Stop at the line, on purpose.** Notice what just happened. My question was about a *seller* and whether shops are linked. It was not about a person's home, their partner, or where they sleep at night. The trail *could* keep going — I could try to pin down the exact address, the partner tagged in that photo, the daily routine. I do not. Not because I could not, but because I *should not*. That material is outside my objective, it is intrusive, and chasing it would turn a fair

inquiry into surveillance of a private life — a real person's, with a real family, who came into this only as the name behind a shop. So I write a note — an honest one: *"Trail extended toward home address and family; stopped here. Out of scope and intrusive. Not collected."* Knowing where to stop is not weakness. It is the whole reason people can trust this craft at all.

The skill that makes you a professional is not how far you can go. It is knowing where to stop — and writing down that you stopped.

Step back now and look at what we actually built — and what we threw away. From one shop name, we reached a first name, two confirmed accounts (after binning a third that was a stranger with the same handle), a public email, a second linked shop, and a phone number that we graded weak but that still tied the two shops together. Likely one operator who goes by Jordan — with a surname that surfaced once and stays uncorroborated. Each step was a small, checkable pivot. Each finding was weighed on the confidence ladder before it earned the next step — and the rejected account and the "no records" blank are part of the trail, not embarrassments hidden from it. And we drew a clean line where the public interest ended and a private life began.

That is the loop from the very first chapter, run for real:
exposure (public scraps left online) → pivot (turn each scrap

into the next door) → **attribution** (carefully deciding it is one person) — all inside the rules.

The Report

A finding nobody can read is a finding that does not count. So we write it up. Short, plain, and honest. A good beginner report answers four things: what I was asked, what I found, how sure I am, and what I did *not* do. Here is ours.

Case: pinewicketcrafts non-delivery inquiry > Date: (today's date) > Objective: Determine whether a real, findable person is behind the shop "pinewicketcrafts," and whether it links to other seller accounts.

What I found. The shop "pinewicketcrafts" appears to be run by a real individual who goes by the first name Jordan. (A surname, "Alvarez," surfaced from a single weak phone lead; I flag it as unconfirmed and do *not* treat it as established.) The seller is linked — at the **strong** rung — to a second, now-closed marketplace shop, "oakwicketgoods," and, at the **moderate** rung, to a woodworking-hobby profile. The links rest on a shared public email address (jordan.pinewicket@example-mail.com), matching product photos, and reused contact details.

What I ruled out. A photo-sharing account with the identical handle "pinewicketcrafts" turned out to belong to a *different, unrelated person* (no craft content, matched on name only). I rejected it. Two sites were unreachable during my username

check, so my picture is not complete, and I say so.

How sure I am, and why. The shop-to-shop link is strong; the tie to the woodworking profile is only moderate; and it is *not* confirmed — I want to be clear about that. The conclusion does not rest on any single clue; it rests on a few independent public sources agreeing — matching craft and photos and a reused email across two shops. The phone number I graded weak on purpose: it is a VoIP/app number, so it corroborates the shop link but places nobody anywhere. If the email turned out to be shared by two people, my confidence would drop, so I flag that as the point that would change the picture.

What I did not do, and what I could not confirm. I did not contact the seller or any account. I did not log in anywhere or interact — all collection was passive and from public pages. I stopped when the trail pointed toward a home address and family members, because that is outside this inquiry and intrusive. I could *not* confirm the seller's exact legal identity or intent; "vanished on an order" is what the buyer reports, and I found no independent proof of fraud — only that two accounts are linked and the operator is likely findable. That distinction matters: I am reporting *connections I verified*, not a crime I proved.

Sources. (A short numbered list of every page, with its web address and the time I viewed it — including the searches that found nothing, and the account I checked and rejected.)

Read that report again and notice its tone. It never oversells. It says how sure, not just what. It admits the gaps. It shows the line I refused to cross. A report like this is something you can hand to another person and they can check every step. That is the goal. Not to sound clever — to be *checkable*.

Write findings, not verdicts. Say how sure you are, show your sources, and admit what you did not do.

A second case: the one that doesn't close

I want to show you one more, and it is the most important in the book — because it *fails*. Real work does this far more often than the tidy story above, and a beginner who never sees it will think they did something wrong when it happens to them. They didn't. Sometimes the honest answer is "I can't say," and reaching that answer cleanly is a professional skill, not a failure.

The request: a buyer left a glowing review on a marketplace that reads like a fake, signed only "Rae M." — is it a real customer or the seller praising their own shop? One name fragment, nothing else.

I work it the same disciplined way. A search for the handle *raem* turns up two live accounts that *both* fit. One is a Rae Molina, active in the same craft scene, in roughly the right region — plausible. The other is a Rae Mercer, who has interacted with this exact seller before — also plausible. Two real people. Both could be "Rae M." I try to break the tie. The email pivot

dead-ends: no public address for either. The photo pivot
dead-ends: neither profile has a usable image. A quoted search
for "Rae M." plus the shop name gives me nothing new. I have
two candidates, each sitting at the weak rung, and *nothing that
lifts one above the other.*

So I stop — and I write the honest ending:

*Finding: Could not attribute. Two plausible individuals
("Rae Molina," "Rae Mercer") fit the fragment "Rae
M."; no public evidence separates them. Confidence:
weak, split across two candidates. I cannot say whether
the review is genuine or self-written. Recommend not
acting on this until a stronger identifier appears.*

Read that. It names the dead end plainly. It refuses to pick the
more *convenient* answer just to hand someone a name. That
refusal *is* the finding. An investigator who forces a confident
answer out of weak evidence is not more useful than one who
says "I can't say" — they are more *dangerous*, because someone
might act on the guess.

*"I could not confirm this" is a finding too. Write it
whenever it is true.*

One more note before you go — who the subject is changes
the rules. Both cases above were consumer scenarios: a scam
seller, a suspect review. If instead you are ever doing this to vet a

job candidate or check an employee, stop and remember the sidebar from Chapter 2: the extra laws (FCRA, anti-discrimination rules, biometric consent) switch on. In that world the discipline grows one more rule — you *stop and do not record* things that are none of the job's business (someone's religion, health, pregnancy) exactly the way we stopped at Jordan's home address here. And for formal hiring decisions, that work often has to run through a licensed screener, not a hobbyist. Different subject, extra rules. Know which world you are in before you start.

Where to Keep Learning

You have finished the foundations. Now the only way forward is reps. Here is how to get them safely and for free.

Practise on yourself first. You are the one subject you are always allowed to investigate fully. Search your own name, your own old usernames, your own email in a breach checker. It is eye-opening — and a little uncomfortable — to see your own footprint the way a stranger would. Bonus: everything you find, you can then go *clean up*. Best OPSEC lesson there is.

Do OSINT challenges and CTFs. "CTF" means capture the flag — a game where you solve puzzles to find a hidden answer. There are OSINT-flavoured ones where you geolocate a photo, or trace a fake persona across sites. They are safe (the "targets" are invented for the game), they are graded, and they build the exact

pivot muscles from this chapter. Search for beginner OSINT CTFs and quizzes; many are free and browser-based.

Try a supervised missing-persons event — carefully. Groups like Trace Labs run "search parties" where volunteers use OSINT to help find real missing people, working strictly inside the law and a firm no-contact rule. This is real, meaningful work, but it is *not* a beginner free-for-all: you follow their OPSEC guide, you use a research account, and you never interact with anyone. Go in humble, read their participant guide first, and learn from people who have done it for years.

Read the reputable guides, not random blogs. A few sources are trusted across the whole field — for example, Bellingcat's online investigation toolkit and their beginner verification guides. These are written by people who do this for a living and are careful about the law and the ethics. Bookmark the good ones. Come back to them.

Keep the habits, not just the tools. Tools will change. New sites appear, old ones die, apps get renamed every year. The *process* does not. Tight question, take notes from action one, pivot with discipline, weigh every finding on the confidence ladder, stop at the ethical line, write it up honestly. Learn that, and any new tool is just a new door for the same method.

A Bridge to Volume Two

This book was the foundations: the loop, the rules, precise search, usernames, email, phone, the confidence ladder, and the basics of staying safe and staying ethical.

Volume 2 goes deeper. We will get serious about OPSEC — real research accounts, isolated environments, and controlling what *you* leak while you look. We will move into images and places: reverse image search done properly, and geolocation — figuring out where a photo was taken from what is in the frame. We will build timelines and relationship maps, and learn the structured ways professionals fight their own bias so their conclusions hold up. Everything there stands on what you just practised here.

For now, you have done the thing this book was for. You ran an investigation, from one small clue to an honest report, and — the part that matters most — you knew where to stop. That is the whole craft in miniature. There is a person at the end of every trail you will ever follow. Keep your notes. Stay curious. Stay kind. And I will see you in Volume 2.

— Yehor

Exercise

Run your own tiny case, on yourself, using only the steps in this chapter.

- **Pick your seed.** Choose one old username you used years ago. Write it at the top of a fresh notes file, with today's date

and this objective: *"Find what public accounts and info still trace back to this old username of mine."*

- **Search it exactly.** Put the username in quotes in a search engine. Note every result with its web address and the time you looked.
- **Pivot it across platforms.** Use a free username-checker to see where that name still exists. Note each hit as a *lead*, not proof.
- **Confirm or reject.** For each hit, decide: is this really me, or someone who picked the same name? Write your reasoning and a confidence level for each.
- **Find your own line.** At some point you may find something private — an old address, a family member. Practise stopping. Write the note: *"Reached private info; stopped; out of scope."*
- **Write the four-part report.** What you were asked, what you found, how sure you are, and what you did not do. Keep it under one page.

Then do the useful part: clean up whatever you found. That is your first real investigation *and* your first real OPSEC win.

Recap

- **One clue is enough to start** — a single public scrap, worked with a tight question and disciplined pivots, can carry you all the way to an answer.

- The mess is the job, not a mistake — real trails have wrong hits, dead ends, blank results, and tools that time out. A run with no rejected leads is not a clean run; it is a careless one. Write down what you threw out and why.
- Weigh before you walk — every finding gets a rung on the ladder (weak → moderate → strong → confirmed), and you never build the next step on an unconfirmed lead; only several independent sources agreeing earns confirmed.
- The report and the line are the craft — a professional writes findings honestly (how sure, what sources, what they did *not* do) and knows exactly where the public interest ends and a private life begins.

Epilogue

What you can do now — and what comes next.

Look back at where you started this book.

A username was just a username. An email was just an address. A phone number was just digits. Now you know the truth: each of them is a thread, and you know how to pull it — gently, in order, confirming as you go.

You learned the one loop the whole craft runs on — exposure → pivot → attribution. You learned the rules that keep you a professional instead of a problem. You learned to protect yourself before you look at anyone else. You learned to search like someone who knows what they're doing, and to turn a username, an email, or a phone number into the beginning of a real, honest picture.

That is the foundation. Everything ahead is built on it.

The honest part

You will be tempted, now that you can see, to see everything. Resist it. The skill you have started to build is quiet power, and quiet power is judged entirely by how you use it. There is a real person at the end of every thread — someone who never agreed to be looked at. Look only when you have a good reason. Confirm before you conclude. And when the trail leads

somewhere you have no business going — stop. That instinct, more than any technique in this book, is what makes someone trustworthy with this work.

Where you go from here

The next thing to do is simple: practice. Run the loop on yourself. Run it on things you own. When you're ready for real, guided practice, the OSINT community runs open challenges and events — safe, legal places to sharpen everything you just learned. (You'll find a starting list in the appendix.)

Then, when the single clues start to feel easy, you're ready for Volume 2 — Cross-Platform Attribution, where one thread becomes many, and you learn to follow a person across every place they exist at once.

You came in able to search. You're leaving able to *investigate*.

Every shadow has a name. Now you know how to begin reading one.

— Yehor Selin

Appendix

A glossary, the method on one page, and where to practise.

A — Plain-words glossary

- **OSINT** — finding things out using only public, open information.
- **Exposure** — a trace someone left in the open (a username, photo, number, post).
- **Pivot** — using one clue to find the next one.
- **Attribution** — being confident that the clues point to one real identity.
- **Sock puppet** — a clean, separate account used for research, never your real one.
- **OPSEC** — protecting yourself while you work, so your looking leaves no trace.
- **Dork** — a precise search using operators (like quotes or site:) to force better results.
- **Metadata** — the hidden data inside a file (for example, when and sometimes where a photo was taken).
- **Confidence ladder** — how sure you are, in four fixed words used all through this book: **weak** (one unconfirmed clue), **moderate** (two independent clues agree), **strong** (three

or more agree, nothing contradicts), confirmed (independently nailed down and checkable).

B — The method on one page

1. Start from one clue. Write down exactly what you have. 2. Ask a small question. "What can this one thing give me?" 3. Pivot. Turn each answer into the next question. 4. Climb the confidence ladder. Never conclude on one clue — confirm with two or three independent ones. 5. Write it down as you go. Save what you found and where you found it — screenshot it and save a copy to a web archive before the page can change. 6. Know when to stop. No good reason, or crossing a line? Stop. And if the evidence never gets past weak, "could not confirm" is a real, honest answer.

Exposure → pivot → attribution. Confirm before you conclude.

C — Where to practise (safely and legally)

- Investigate yourself first. It's the one target you're always allowed to fully explore.
- OSINT challenges and CTFs. The community runs open, legal puzzle events built exactly for learning these skills.
- Missing-persons search events (such as those run by Trace Labs) let you practise for a genuinely good cause, alongside experienced people.

- Public investigation guides from respected teams (for example, Bellingcat's toolkit and exercises) give you real problems to solve.

Start small, stay legal, and keep the method in front of you. Skill comes from reps, not from tools.

A note on the tools in the exercises. They come in two shapes. Some are ordinary websites — you just open them in a browser and type. Others (many username checkers and email tools) are small free programs you install, and some of those run in a *terminal* — a text window where you type commands — rather than with buttons. Neither is better; they just have different on-ramps. If a tool wants a terminal and that is new to you, do not let it stop you: start with a browser-based equivalent, learn the method, and pick up the terminal ones later. And you never need to spend a cent to practise anything in this book.

D — Take it back: shrinking your own footprint

Once you have investigated yourself, you have a map of your own exposure. Now use it. Every technique in this book runs on scraps you left in the open — so the defence is simple: leave fewer scraps, and take back the ones you can. Work down this list in order; the first items give you the most protection for the least effort.

- Change the passwords that leaked. Run your emails through a breach checker (Chapter 6). For every address that

shows up in a breach, change that password *and* every other place you reused it. This is the single highest-value thing on the page.

- **Turn on two-factor everywhere that matters.** Email first, then anything with money or messages. An app-based code beats a texted code. A leaked password alone stops being enough to get in.
- **Break the username thread.** You saw in Chapter 5 how one reused handle links all your accounts. For anything you want kept separate, use a different name. Retire old handles you no longer use.
- **Close and delete dead accounts.** The forum you used once in 2016 is still leaking. If you will not use it again, delete it — do not just abandon it. Fewer doors, fewer pivots.
- **Turn off "find me by my number / email."** On social apps and messengers, switch off the settings that let a stranger find your profile from your phone number or email. This kills the phone-pivot from Chapter 7.
- **Uninstall contact-name apps and revoke contact access.** Those caller-ID apps uploaded your whole address book — and made your number searchable to others. Remove them, and revoke contacts permission from apps that do not need it.
- **Strip location and metadata before you post.** Photos can carry the place they were taken. Turn off location tags in your camera, and check the privacy settings that geotag your posts.

- Lock down what strangers can see. Set old posts and photo albums to friends-only. Review the "public" side of each profile as if you were a stranger looking in — because someone will be.
- Use the removal routes that exist. Search your name; where a data-broker or "people-finder" site lists you, most have an opt-out or removal form. It is tedious, but it works. Re-check every few months — some relist you.
- Recheck yourself on a schedule. Exposure grows back. Run the self-investigation from Chapter 8 twice a year and clean up whatever is new.

When the person looking already knows you. This whole book leans on one comfort — that a stranger has to *guess* whether two accounts are the same person, and guesses can be wrong. That comfort disappears when the searcher is someone who already knows your old usernames, your number, your email, your face, and your family's names — an ex, a stalker, a harasser. For them, the "same name isn't the same person" protection does not apply; they have the ground truth. If that is your situation, understand two hard things. First, a real photo you already posted, or a relative who tagged you, can sometimes not be fully un-leaked — the honest fix may be changing the username, number, or email itself, not just hiding it. Second, this is bigger than a checklist. Reach specialist help: a domestic-violence hotline, an address-confidentiality program, or local police —

people who do exactly this for a living.

Protecting your family. Every pivot in this book works on your kids and partner too. Lock down minors' accounts hard: private by default, no location tags, no findable-by-number or -email, and talk with them about what a stranger can rebuild from an oversharing post. Do the self-investigation *on your family's public footprint*, gently, and clean up together.

The best protection is the same skill as the best investigation: know exactly what you are leaving in the open — then decide, on purpose, what stays.

E — Templates you can copy

Steal these. Paste them into a note file and fill the blanks. They turn the habits in this book into something you actually run.

The four-part report. Every finding you hand to anyone fits this skeleton:

```
CASE:           (one line – what this is)
DATE:           (when)
OBJECTIVE:      (the single question, in one sentence)

WHAT I FOUND:   (the answer, plainly)
HOW SURE, AND WHY: (rung: weak/moderate/strong/confirmed
                  + the independent sources it rests on)
WHAT I RULED OUT: (leads checked and rejected, and why)
WHAT I DID NOT DO: (limits, the line I stopped at,
                   what I could not confirm)
SOURCES:        (numbered: URL + date/time viewed +
                 archived copy – include the dead ends)
```


The running note (one row per step). Keep this open from the first click:

TIME		ACTION		FACT (what I saw)	
		OPINION (what I think)		SOURCE + ARCHIVE	

The confidence-ladder cheat card. Never write a conclusion without a rung — and here is exactly what earns each one:

- **weak** — one clue, unconfirmed (a matching name, a single crowd label, a VoIP number). A lead, never a conclusion.
- **moderate** — two *independent* sources agree (e.g. content + dates, or two separate apps). "Independent" means they could not have simply copied each other.
- **strong** — three or more independent sources agree and nothing contradicts them.
- **confirmed** — independently nailed down, and another person could re-check every step and reach the same answer.

If you cannot name the rung, you have not finished the finding.